

Reproducible Science of Deep Learning: The Pruning Case Study

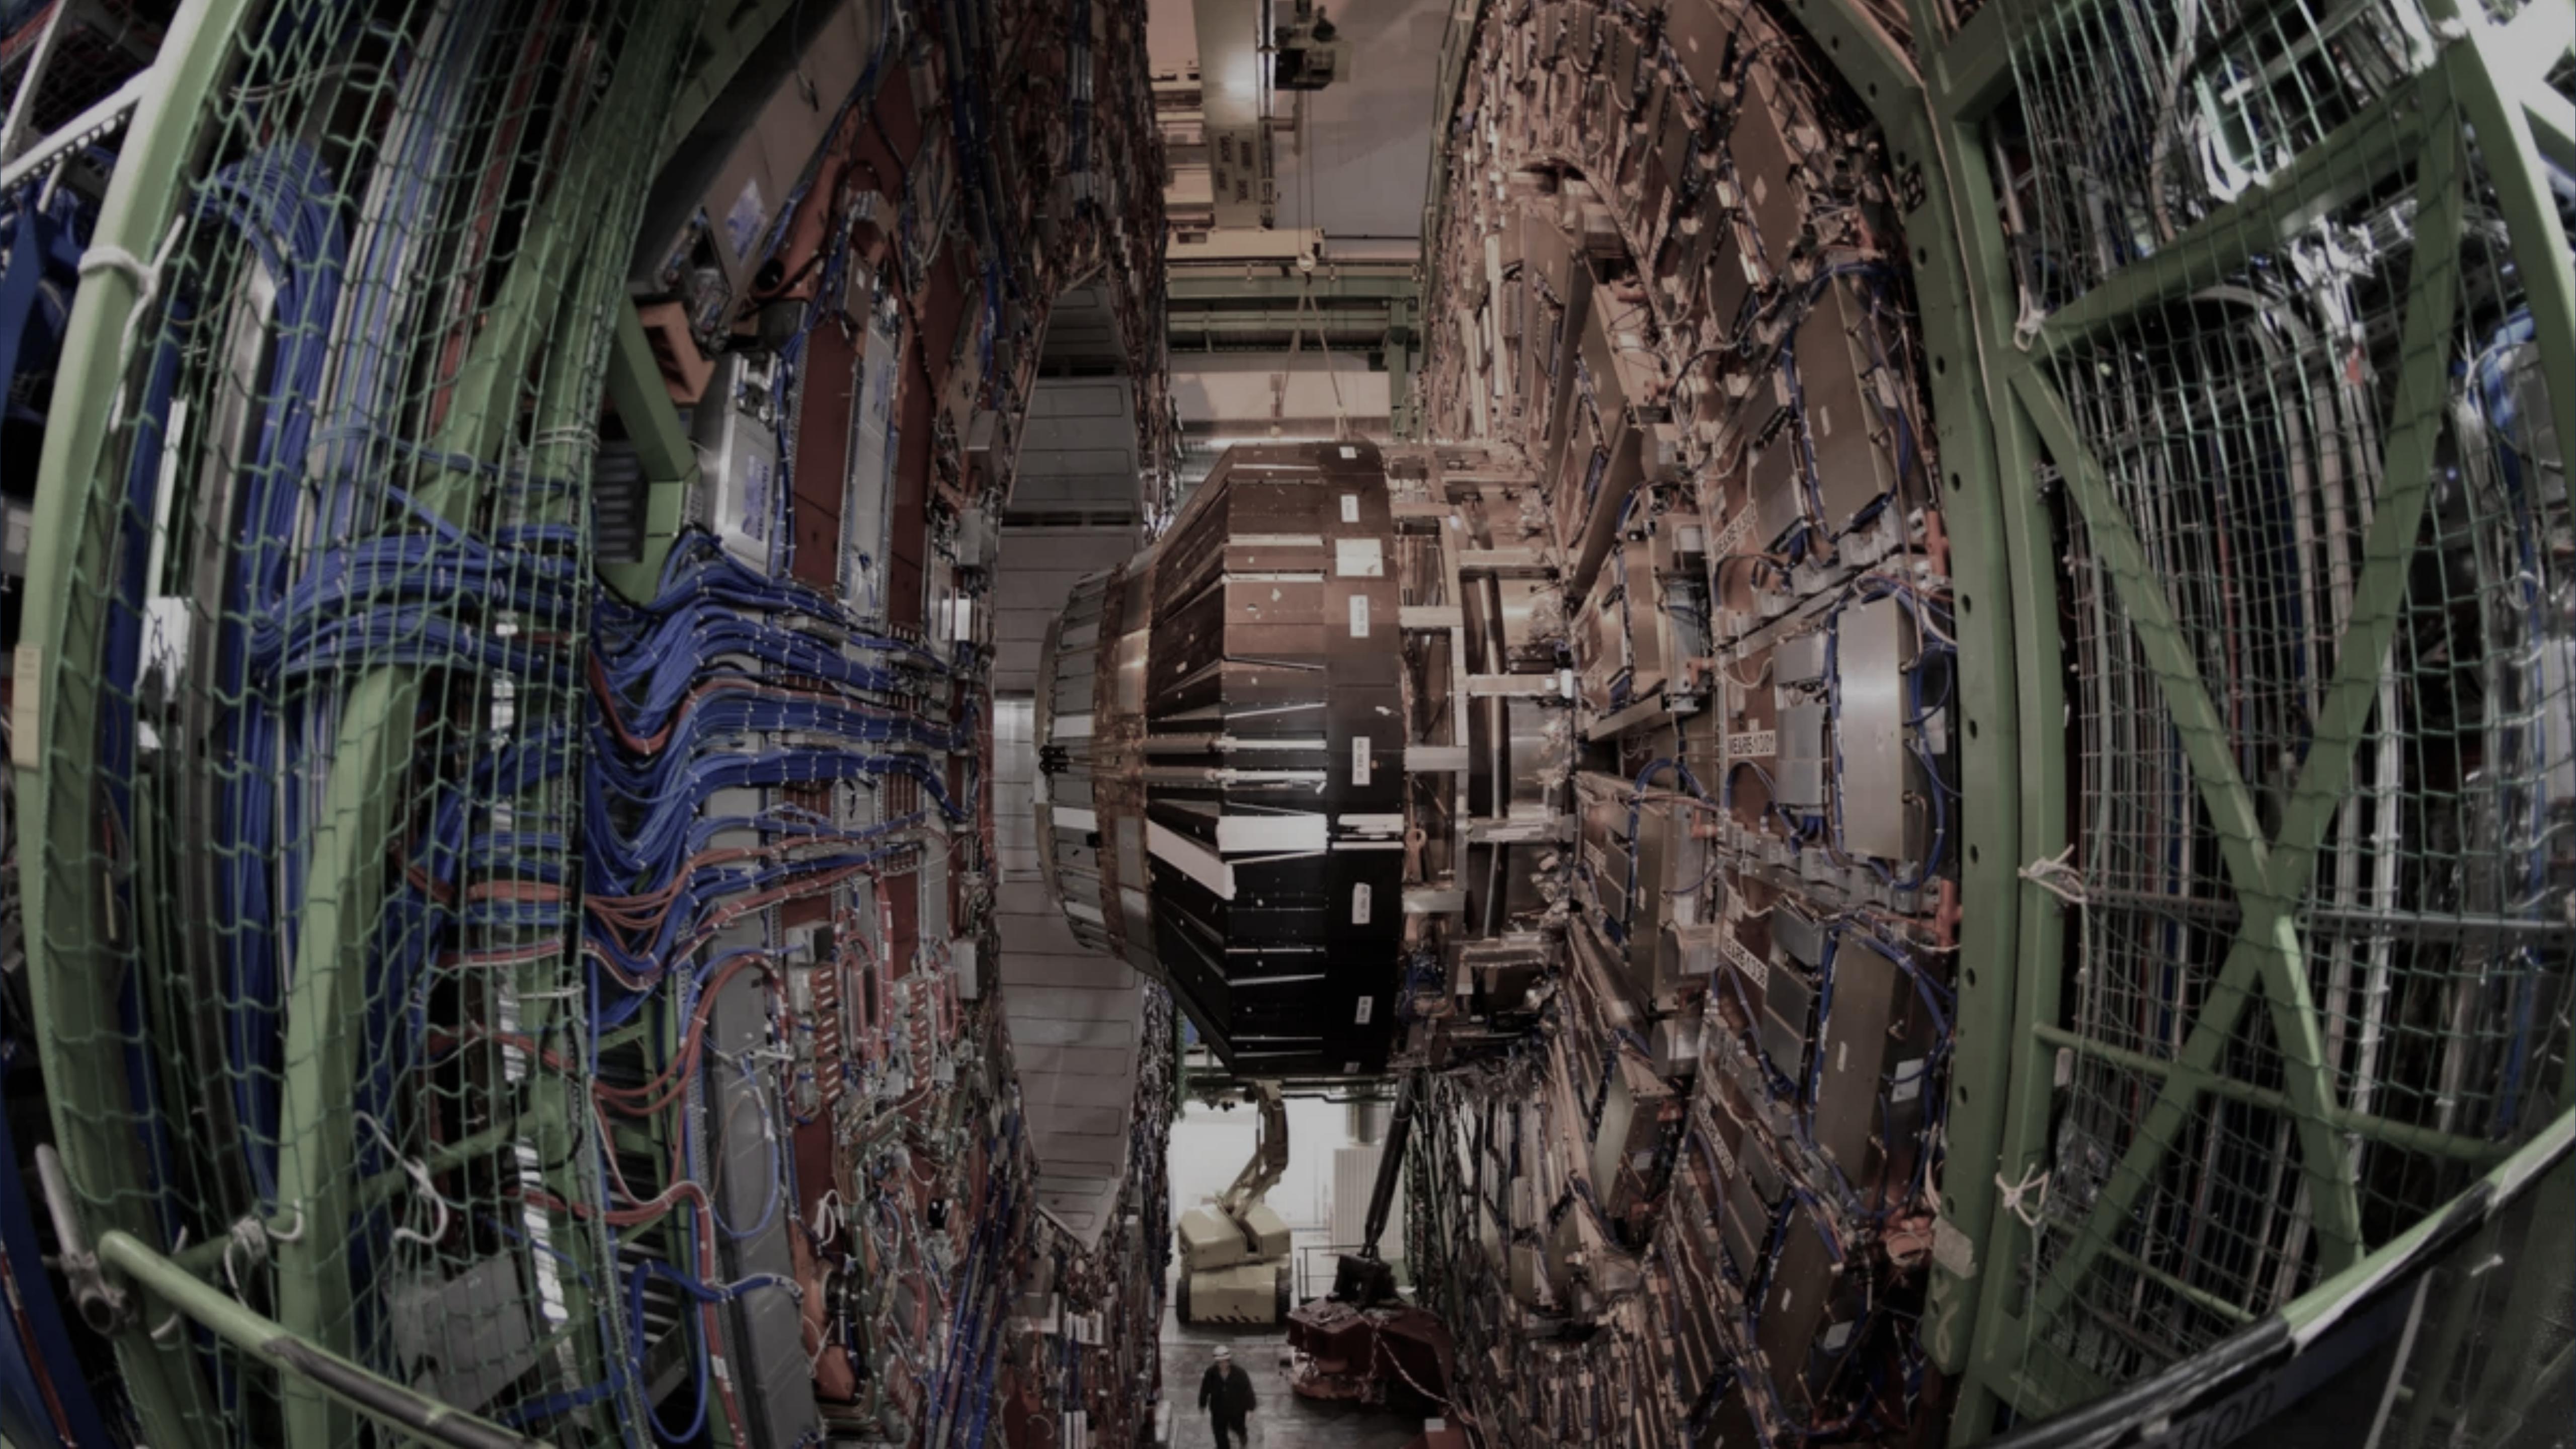
Michela Paganini, Facebook AI Research

 @WonderMicky

UCI Symposium on Reproducibility in Machine Learning

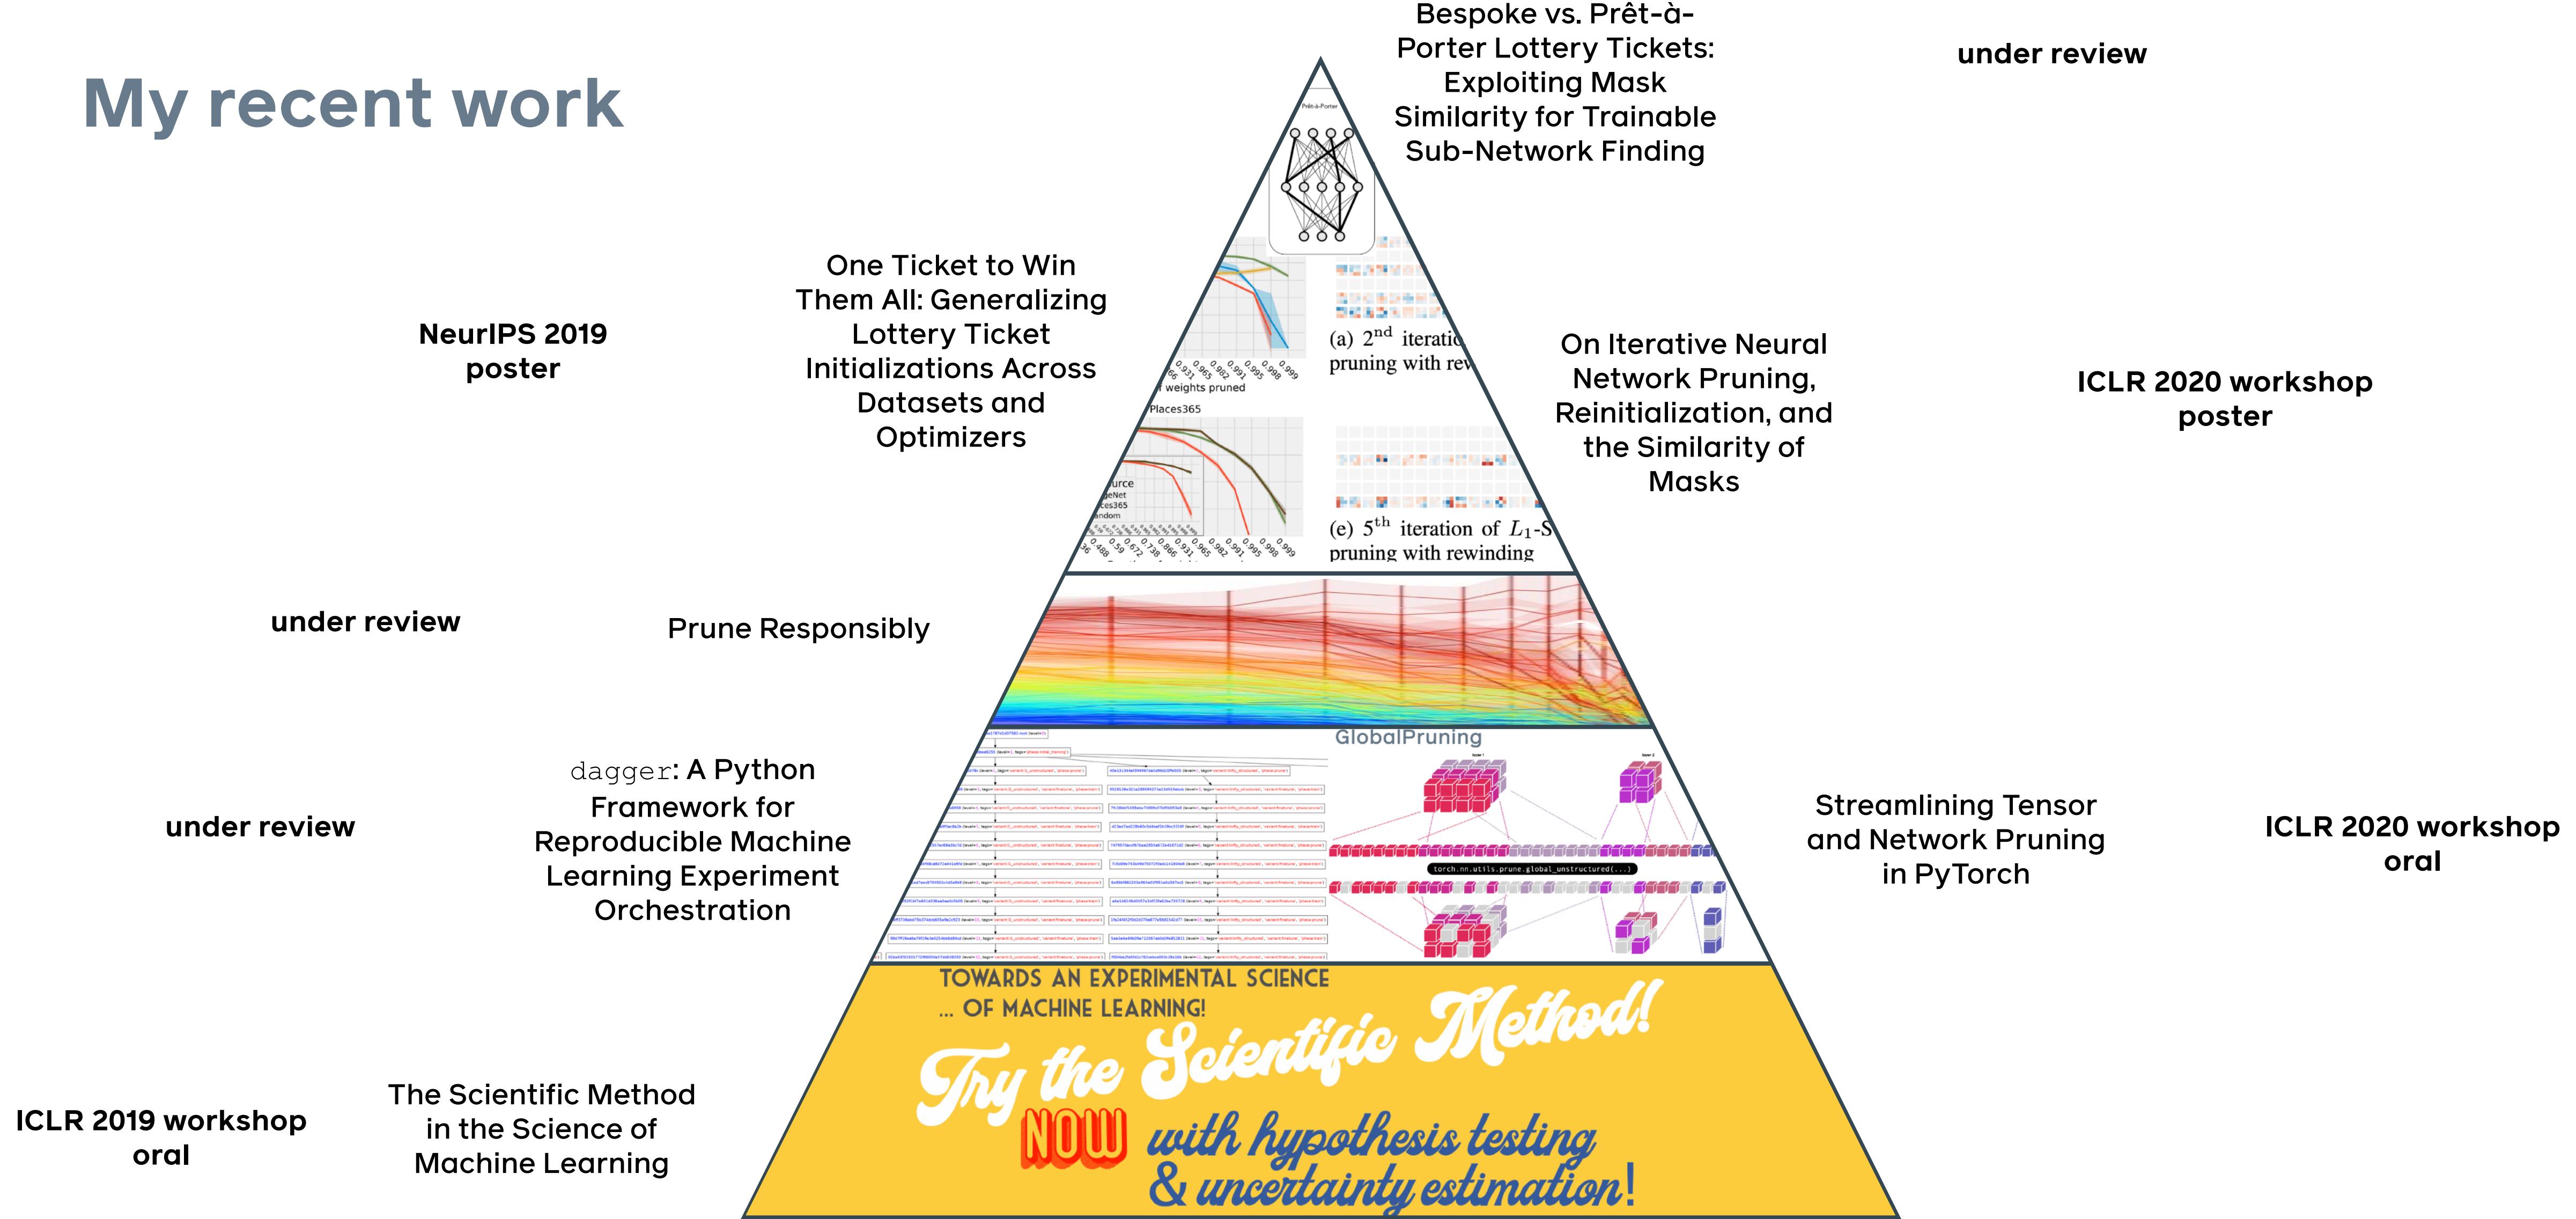
07/22/2020

FACEBOOK



Science is a verb

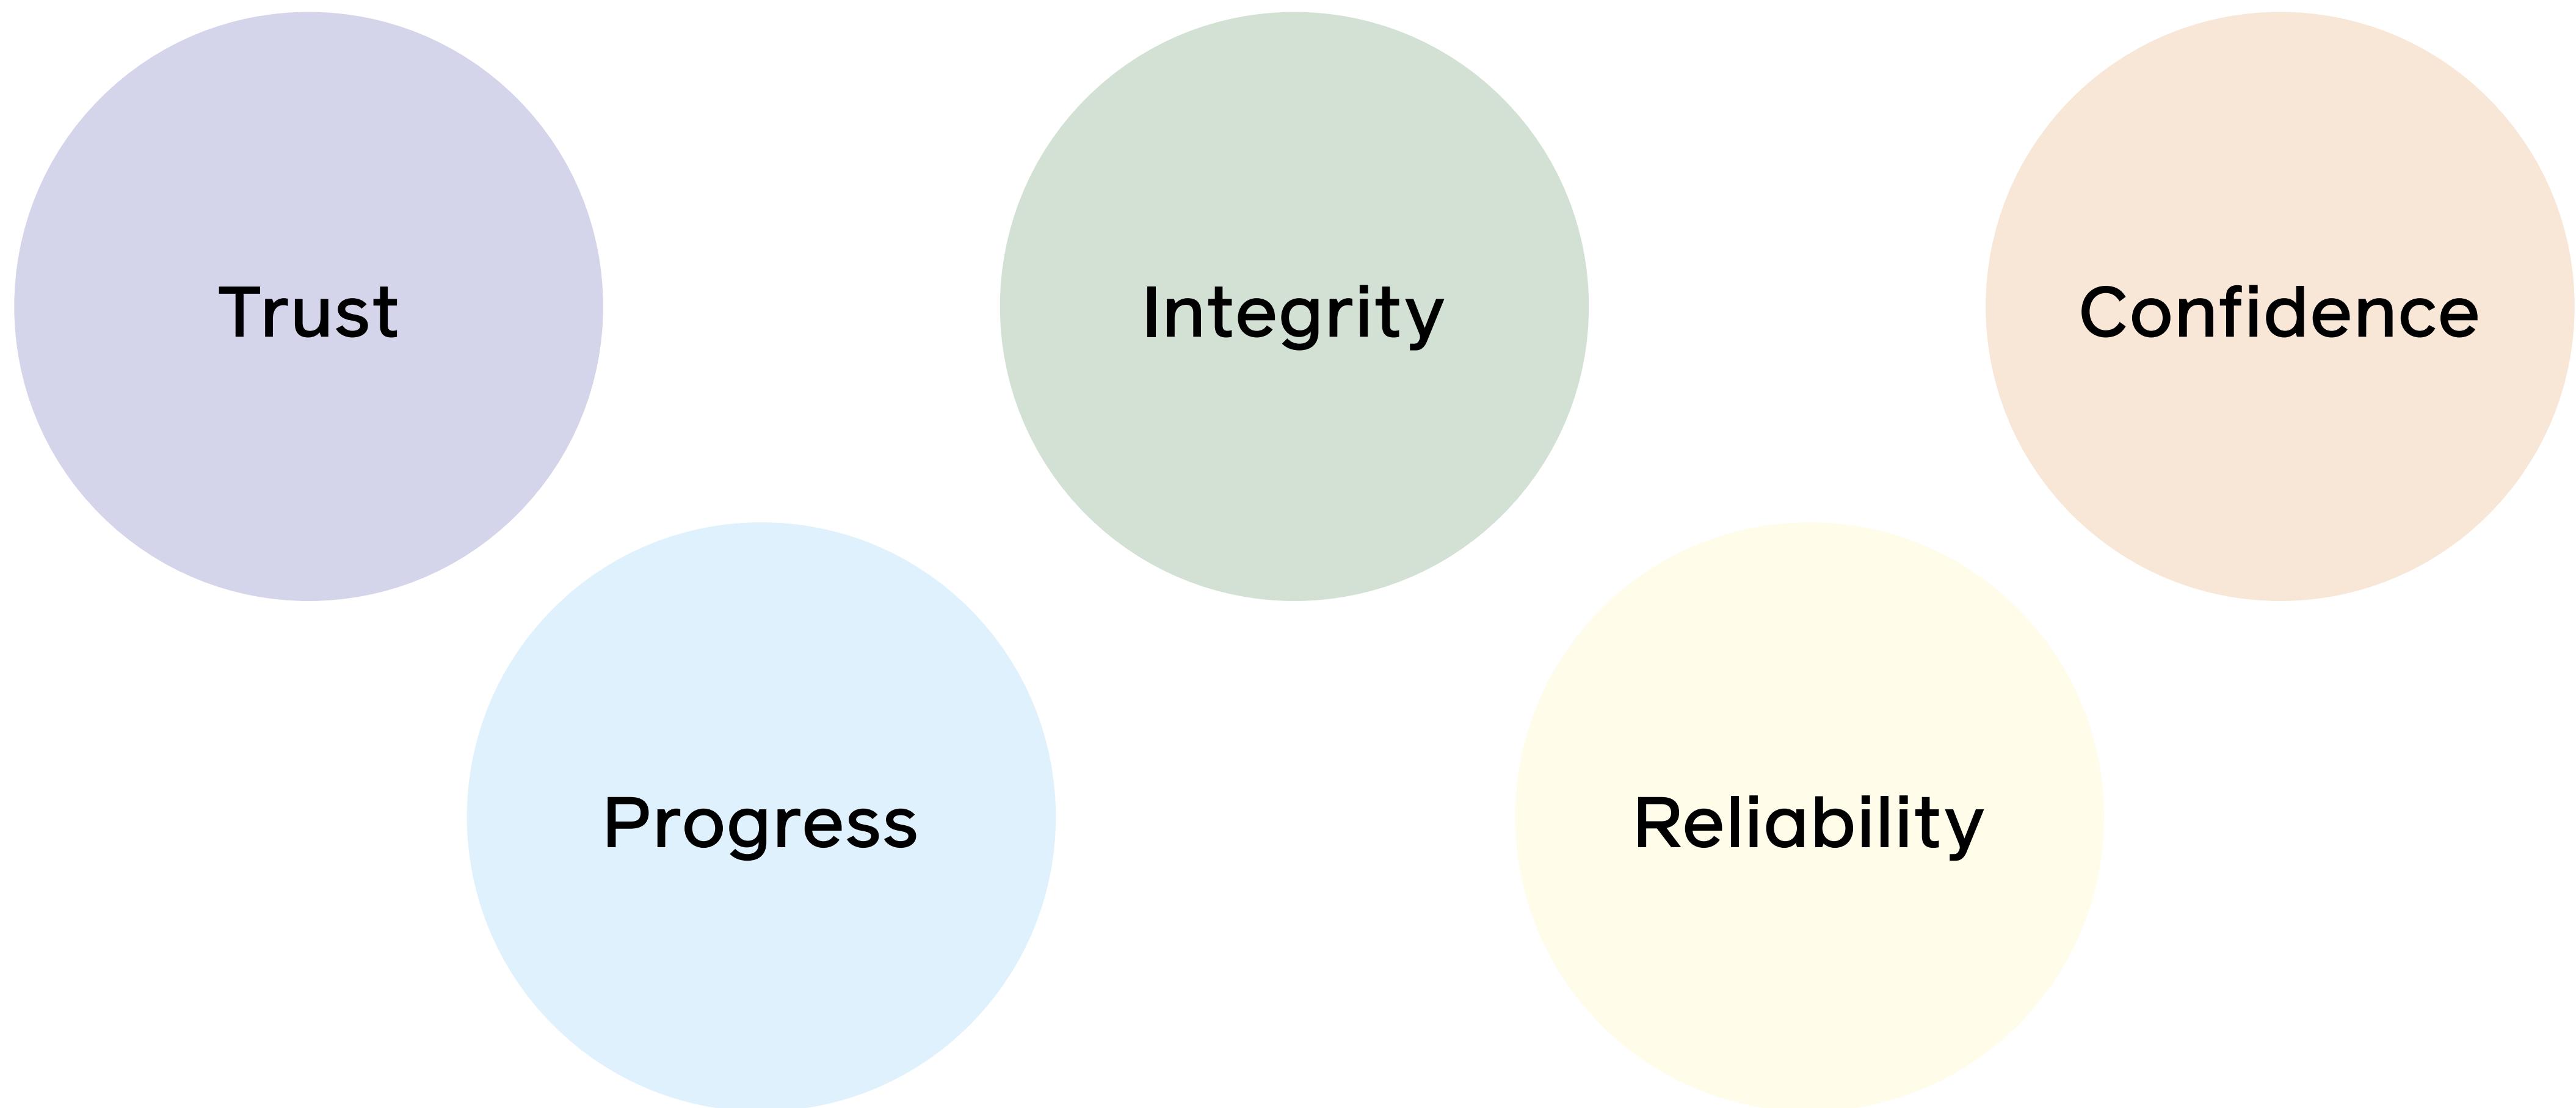
My recent work



Agenda

1. Why reproducibility?
2. The scientific method in the science of ML
3. Pruning for hypothesis testing
4. Dagger
5. Pruning in PyTorch
6. Measuring the disproportionate harm of pruning

1. Why reproducibility?



1. Why reproducibility?

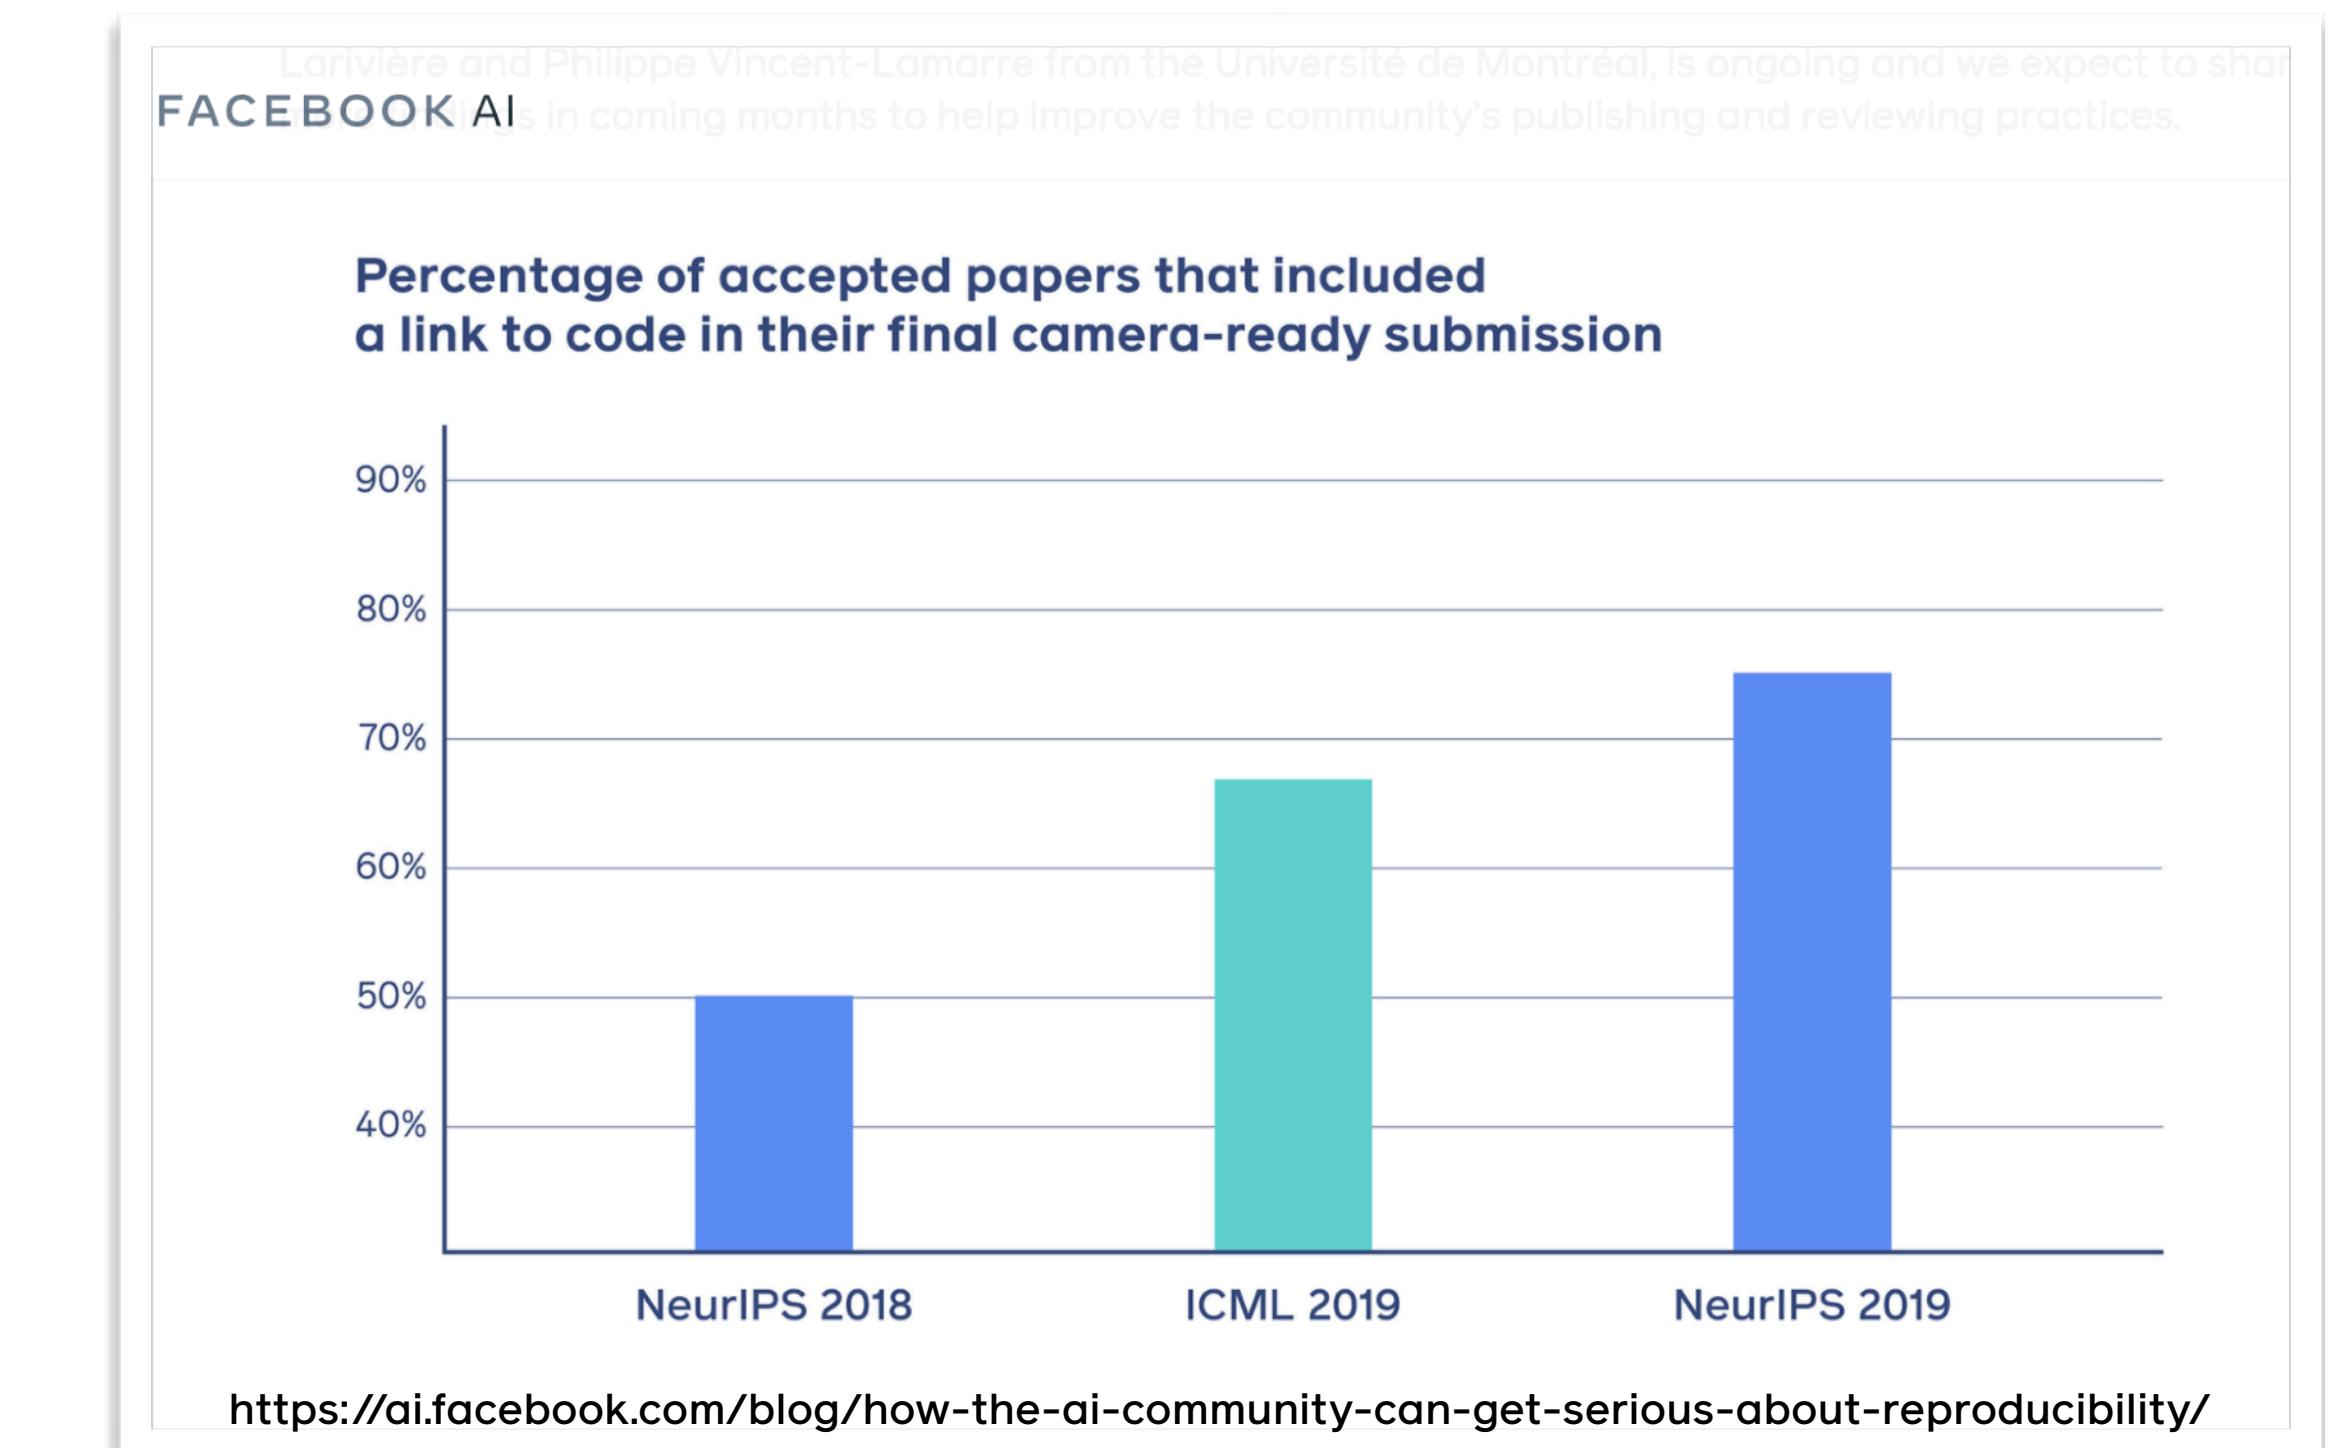
nature > news q&a > article

NEWS Q&A · 19 DECEMBER 2019

This AI researcher is trying to ward off a reproducibility crisis

Joelle Pineau is leading an effort to encourage artificial-intelligence researchers to open up their code.

Elizabeth Gibney





June 10, 2019

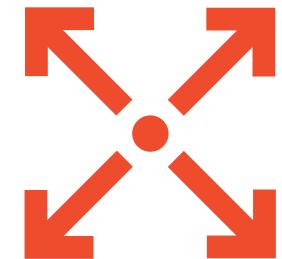
Towards Reproducible Research with PyTorch Hub

PYTORCH HUB

PUBLISHING MODELS

PyTorch Hub supports publishing pre-trained models (model definitions and pre-trained weights) to a GitHub repository by adding a simple `hubconf.py` file.

Discovery



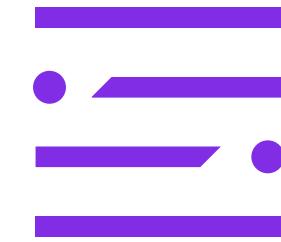
Find the best models related to your research/application!

Reproducibility



Spend minutes instead of days on baselines

Responsibility

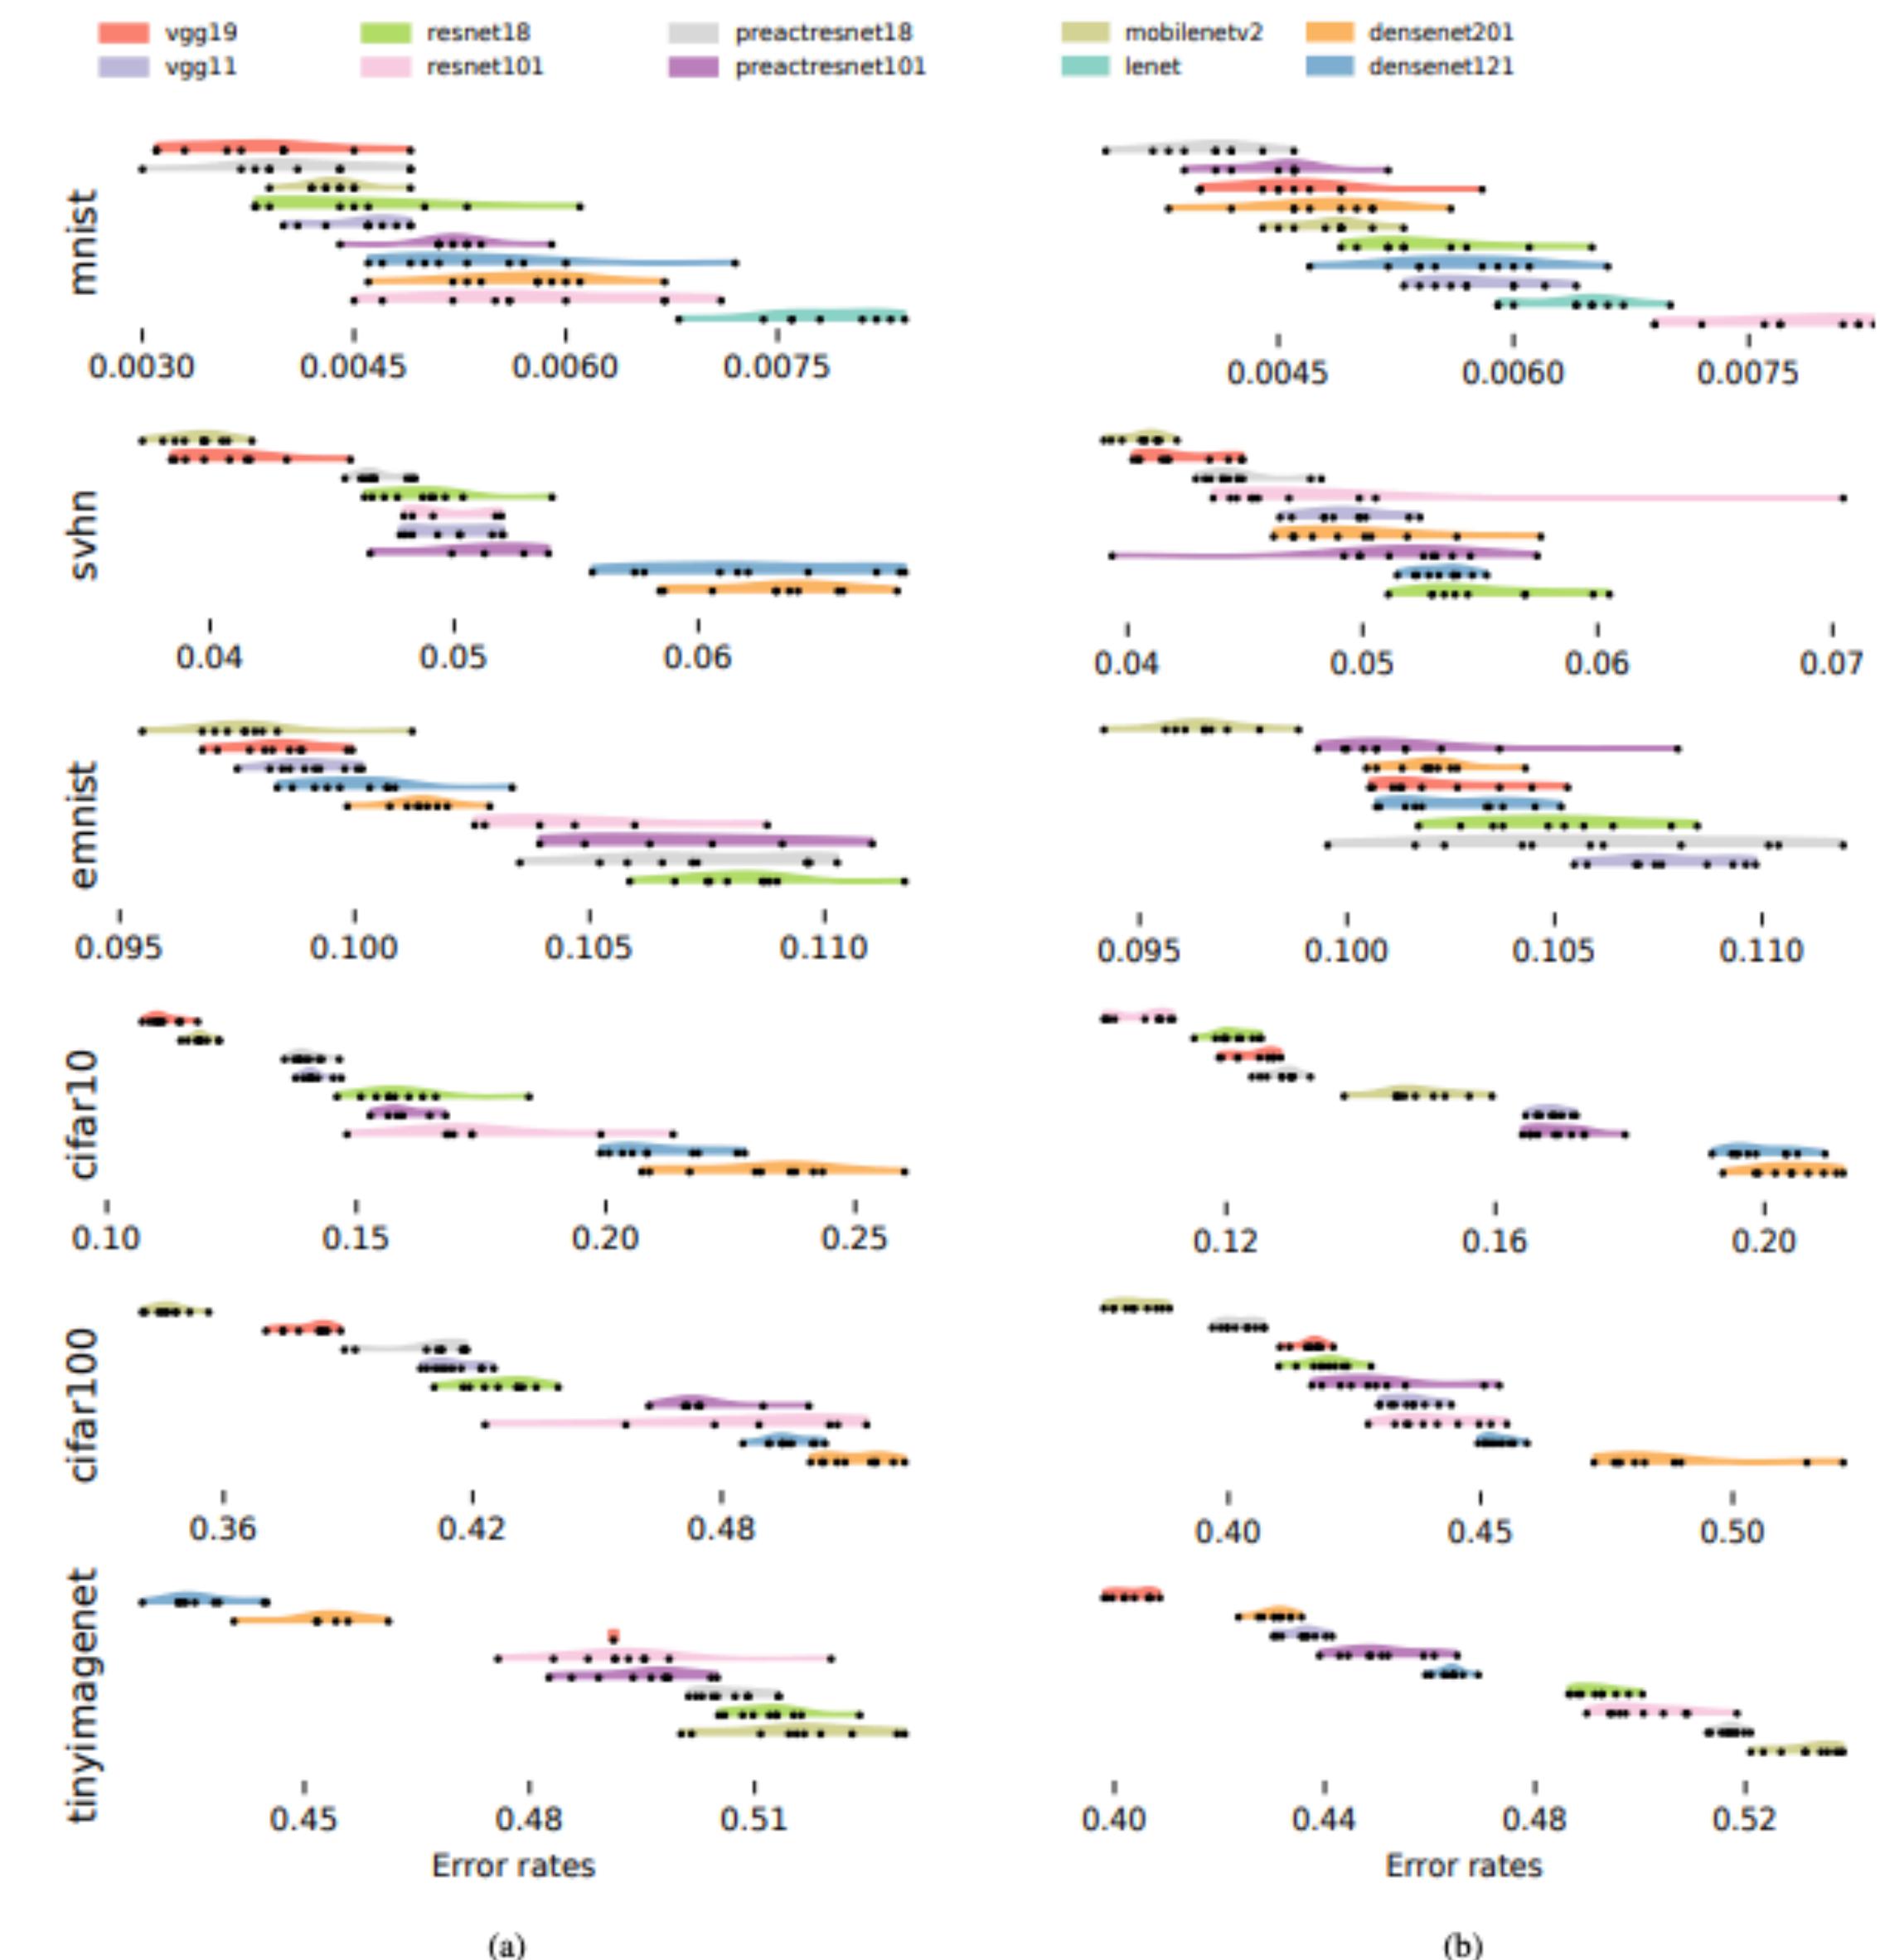


Publish solid papers with reproducible results.

1. Why reproducibility?



- unreproducible findings can be built upon reproducible methods
- not just a matter of deterministic reproducibility of methods and single numerical results
- necessity of ensuring the reproducibility of empirical findings and conclusions by properly accounting for essential sources of variations
- more energy should be devoted to proper empirical research in our community
- promote the use of more rigorous and diversified methodologies



Measurements are affected by sources of variations

What can we learn from the other sciences?

The Scientific Method in the Science of Machine Learning, arXiv:1904.10922

The one and only way to make objective statements?



A social contract among scientists to harmonize workflows and compare findings?



2. The scientific method in the science of ML

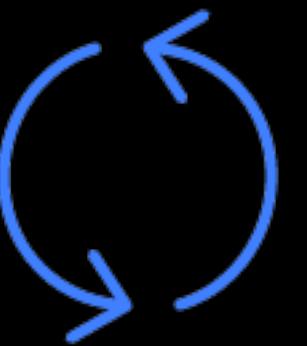
Transparency



Falsifiability



Reproducibility



Intellectual Honesty



Key Steps for Experimental Scientific Research.

hypothesis formulation

statement of expectations

experiment design

statistical analysis

uncertainty estimation

Key Steps for Experimental Scientific Research.

hypothesis formulation

"The null hypothesis is ..., the alternative hypothesis is ..."

statement of expectations

"If the hypothesis is right, then I should expect to observe ..."

experiment design

"I design this experiment to be sensitive to..."

statistical analysis

"Do I observe the expected effect? Is it stronger or weaker than expected?"

uncertainty estimation

"Do I have enough observations and did I account for systematic biases?"

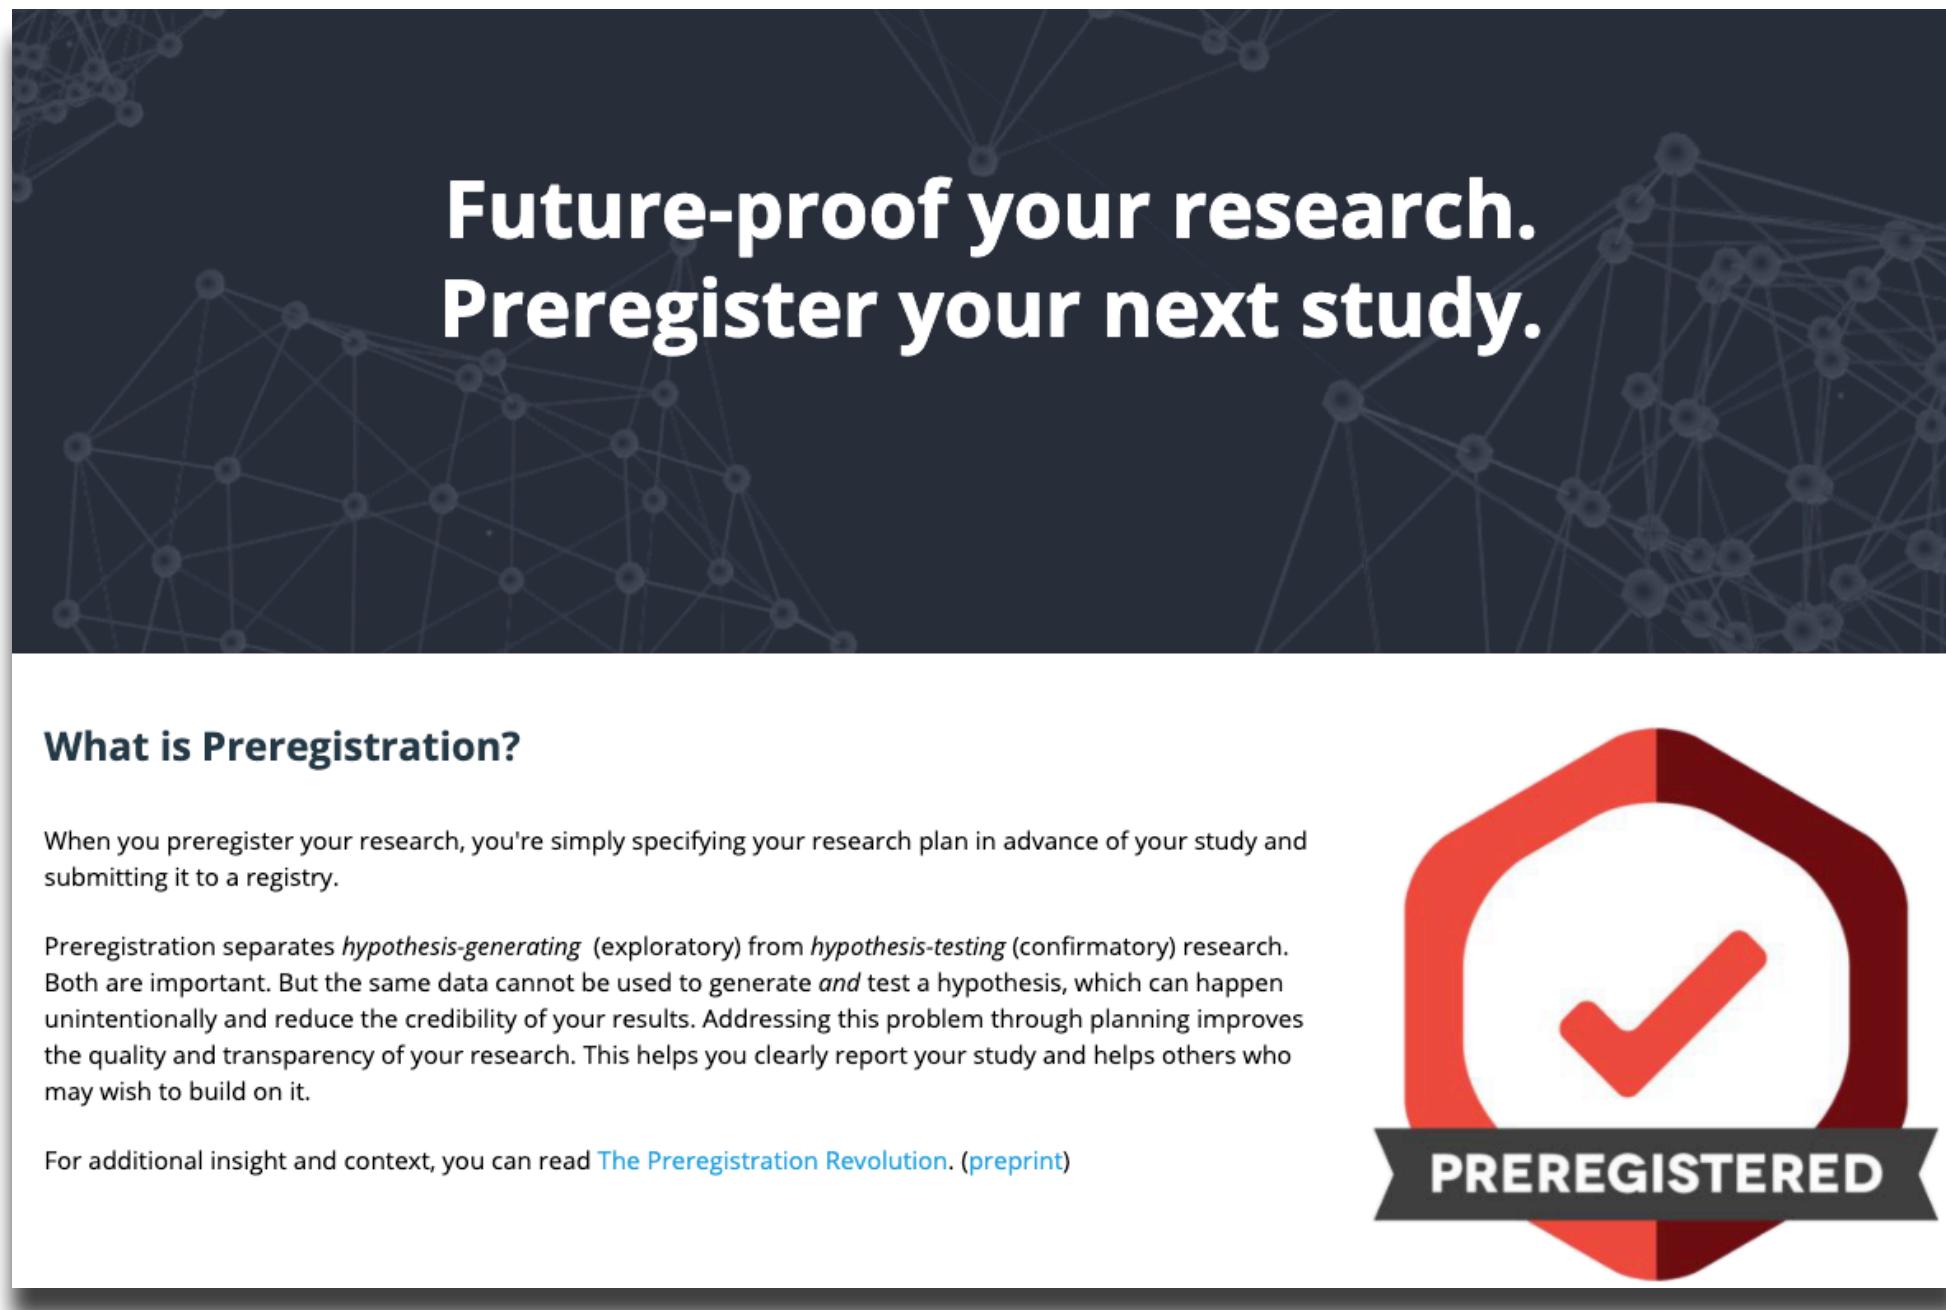
"The first step towards a scientific formulation of ML then demands a more dramatic shift in priorities from drawing and recording single instances of experimental results to collecting enough data to gain an understanding of population statistics."

"it is plausible that a significant percentage of published work claiming state-of-the-art performance actually has no statistical sensitivity to measure their improvement over competing methods."

Blind analysis and pre-registration

Don't judge a paper by its p-value.

cos.io/prereg/

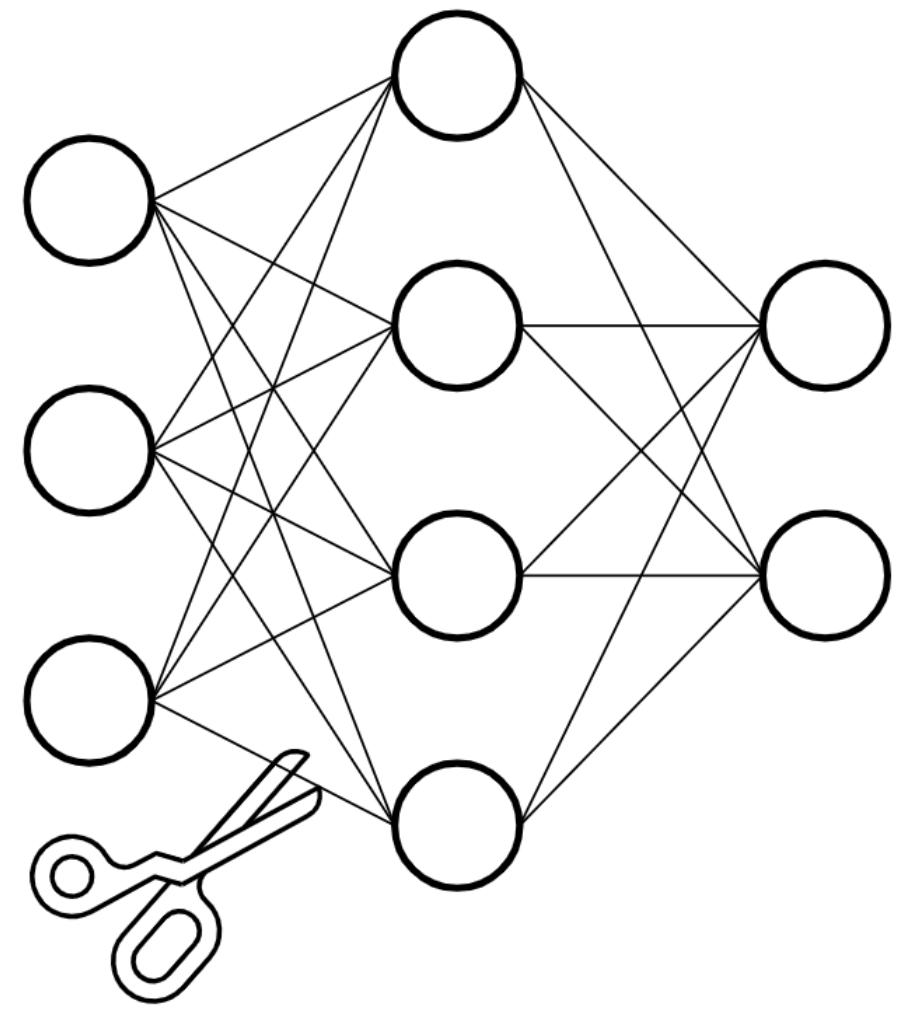


preregister.science

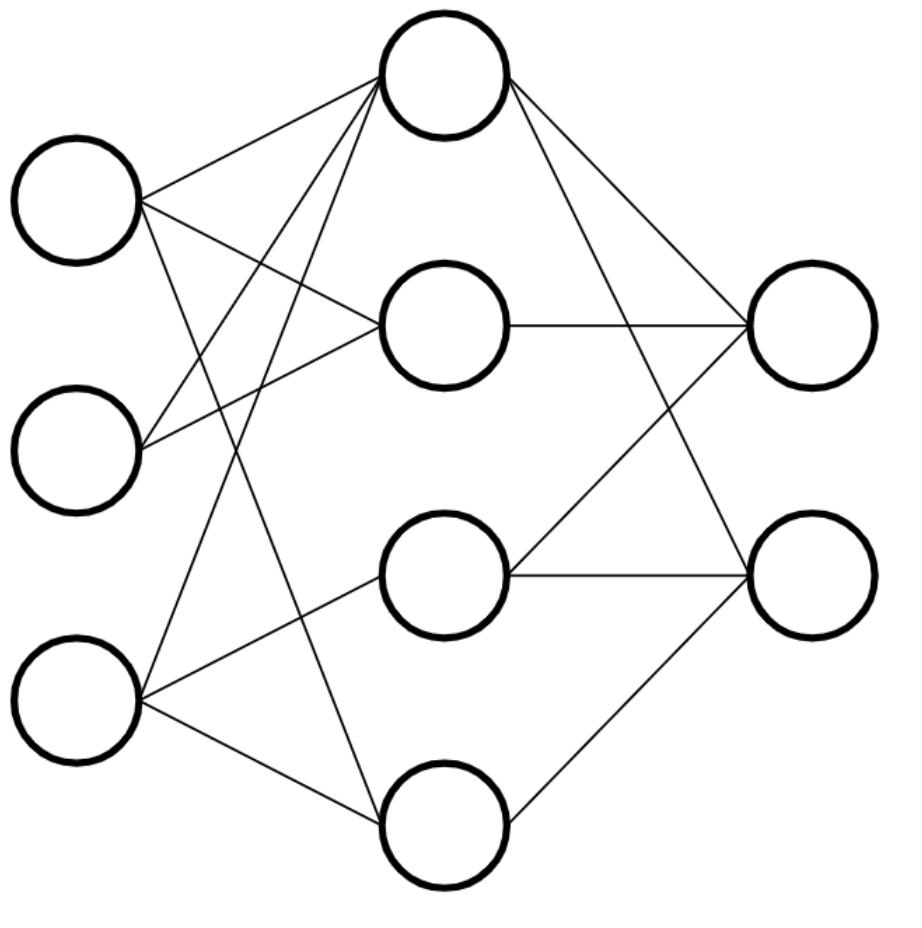


The Pruning Case Study

Pruning



Before pruning



After pruning

"removing superfluous structure"

how to identify?



what kind of structure?



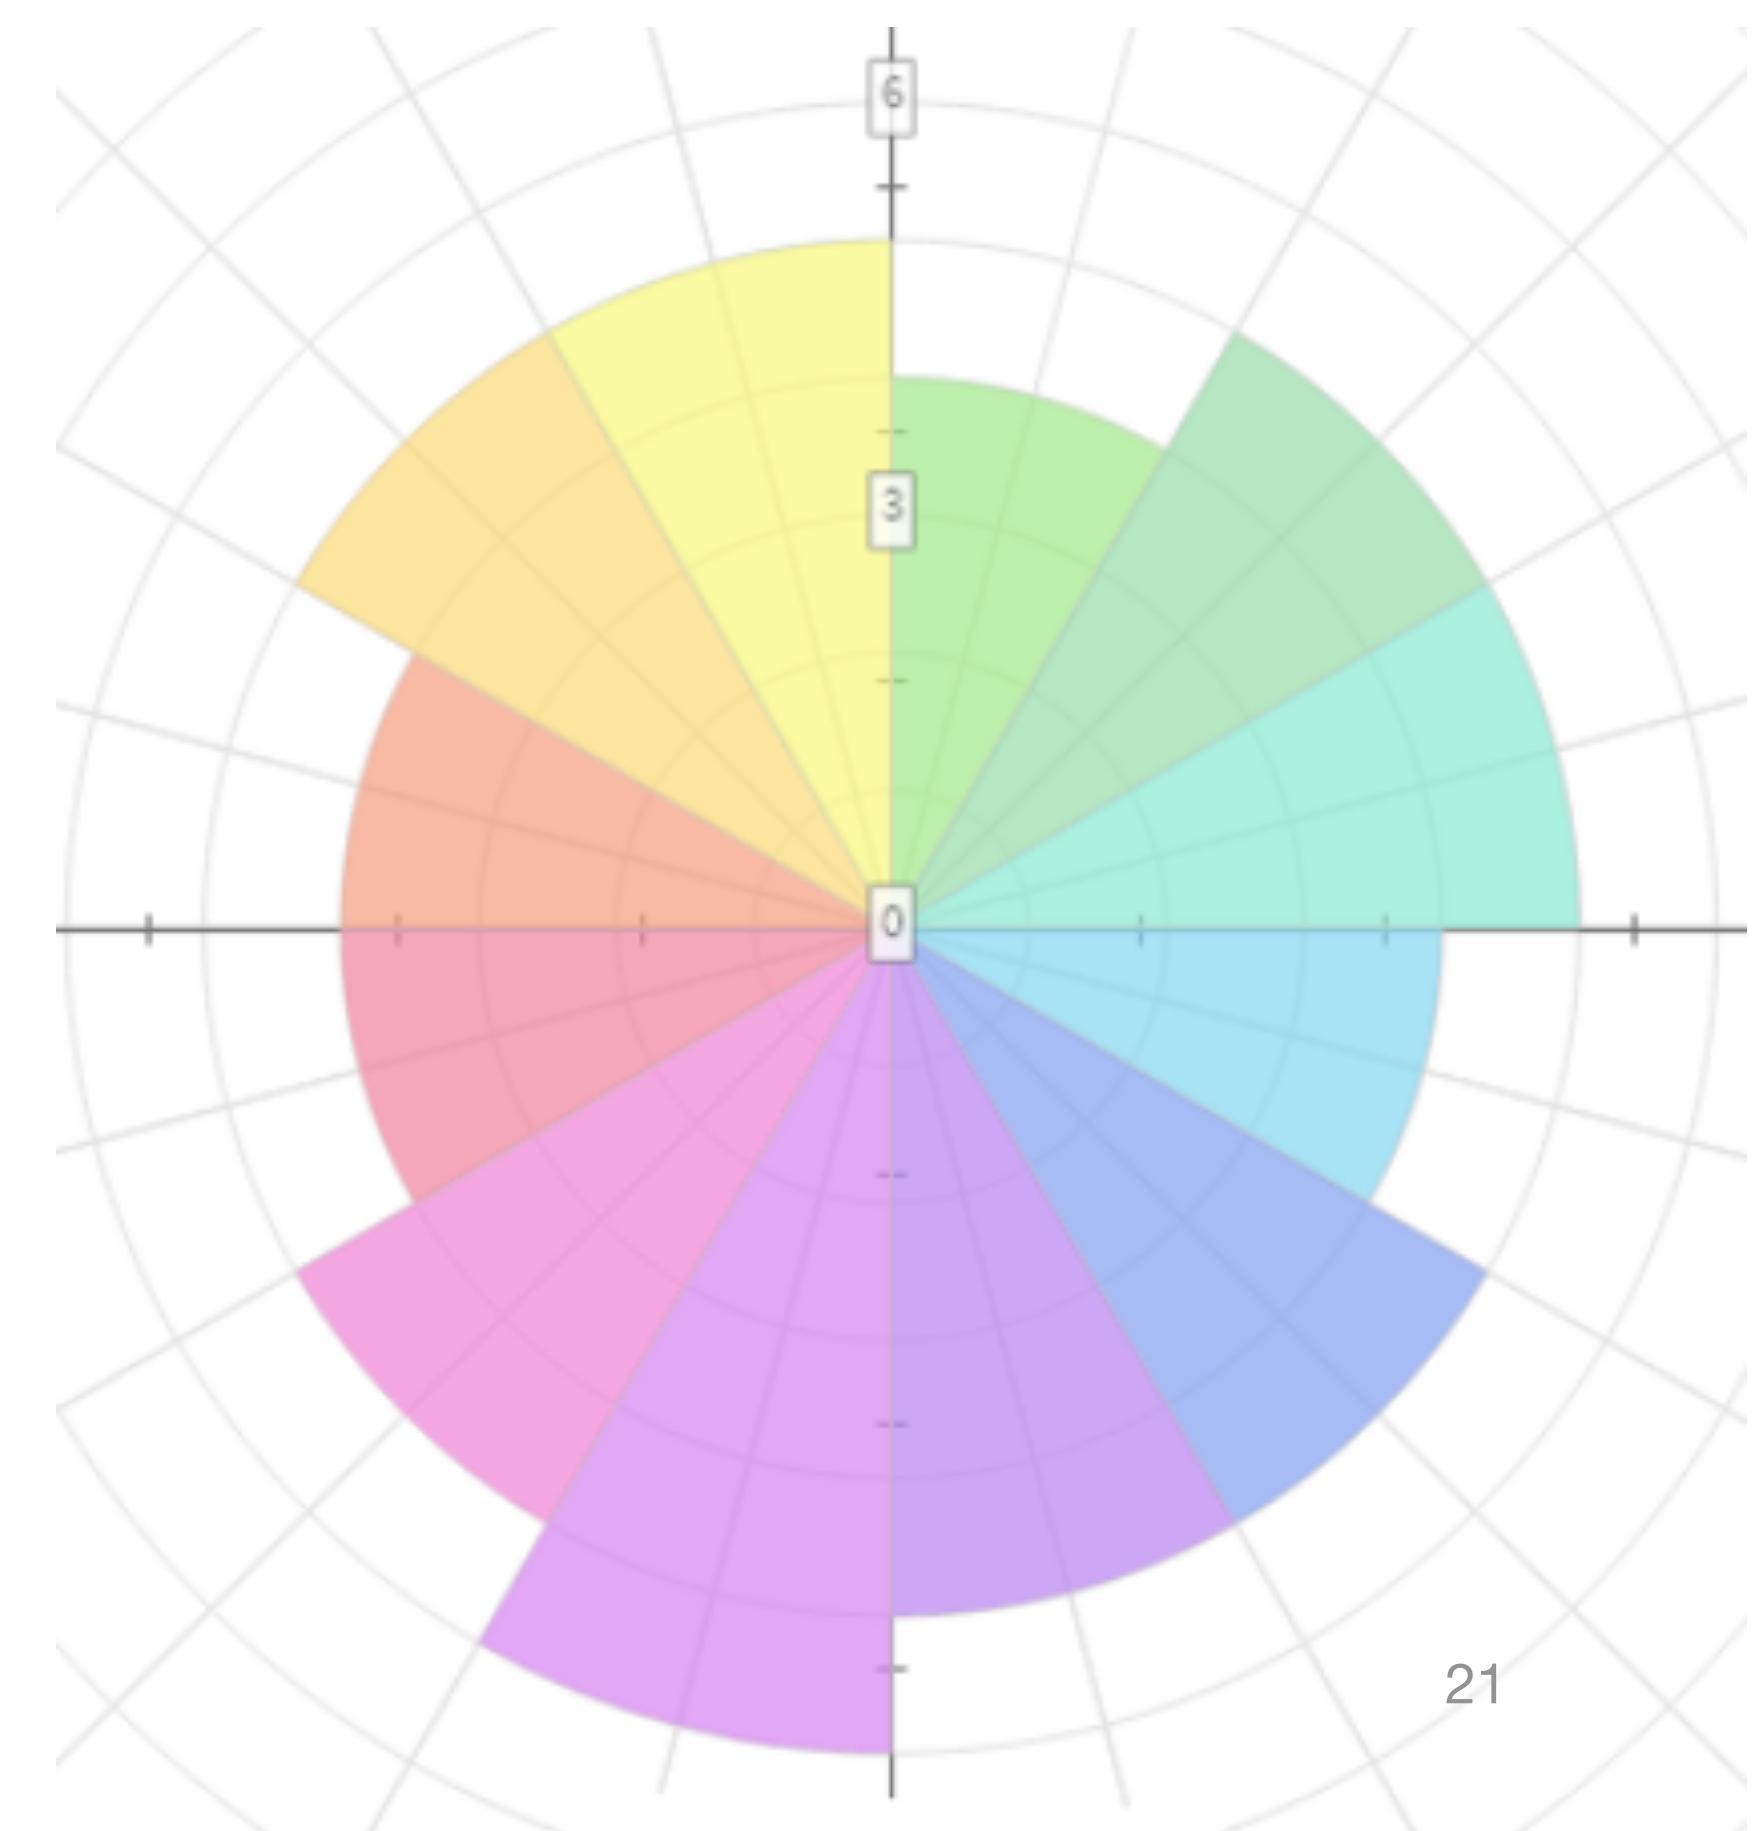
The state of pruning

Pruning should remove unnecessary redundancy and unused capacity

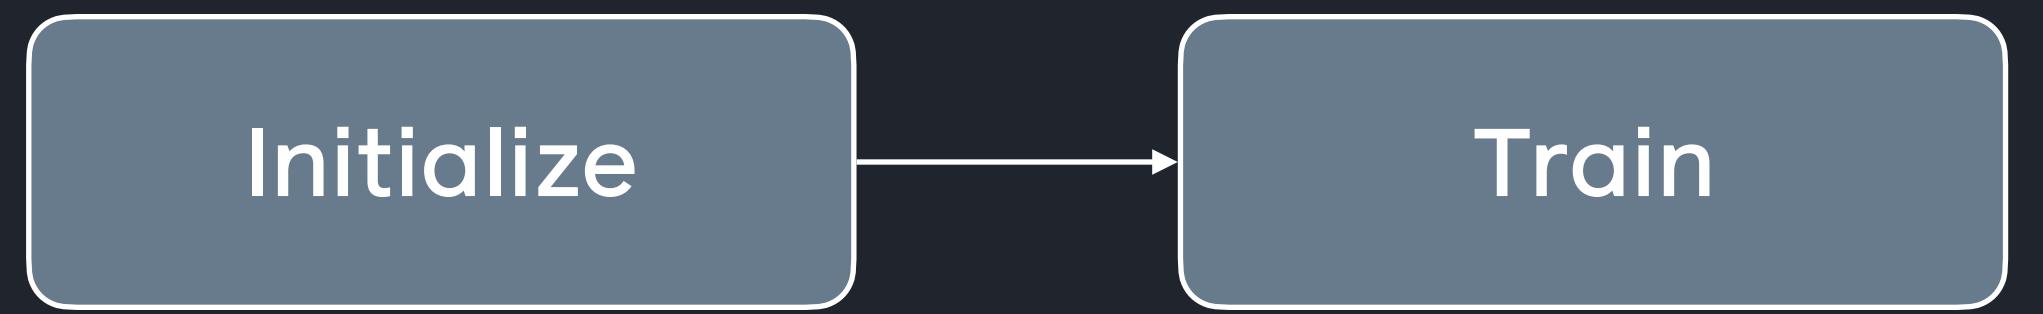
Can be executed *before*, *during*, and *after* training

Pruning methods differ across many dimensions:

- ▶ based on weight magnitude, activations, gradients, Hessian, interpretability measures, credit assignment, random, etc.
- ▶ Layer-wise vs global, unstructured vs structured, etc.
- ▶ Rule-based, bayesian, differentiable, soft approaches, etc.
- ▶ One-shot vs iterative pruning
- ▶ Followed by: finetuning, reinitialization, rewinding

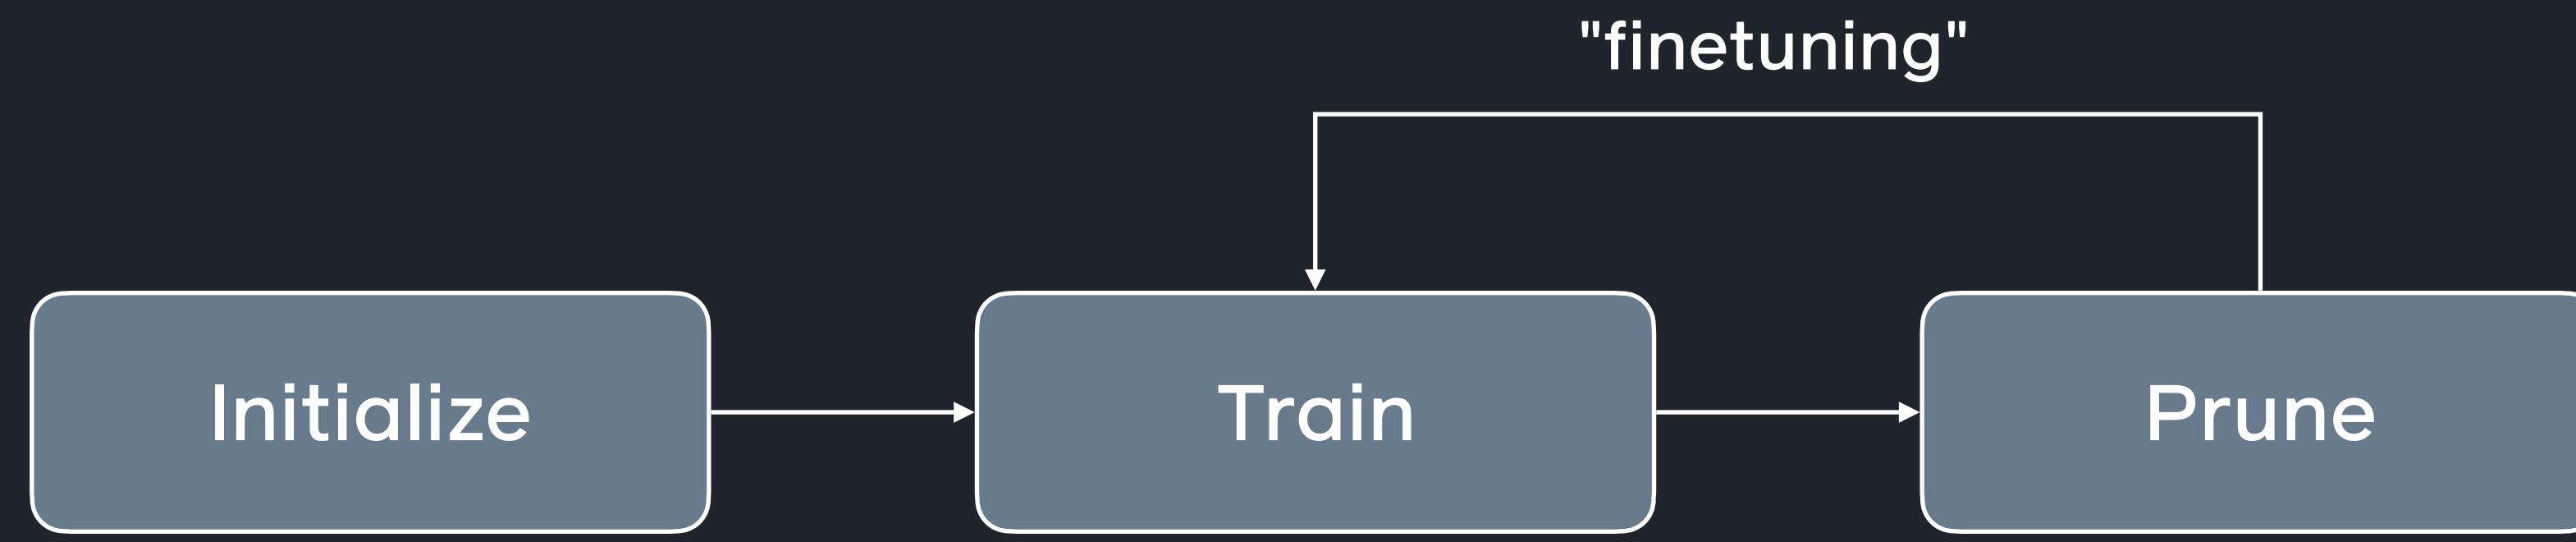


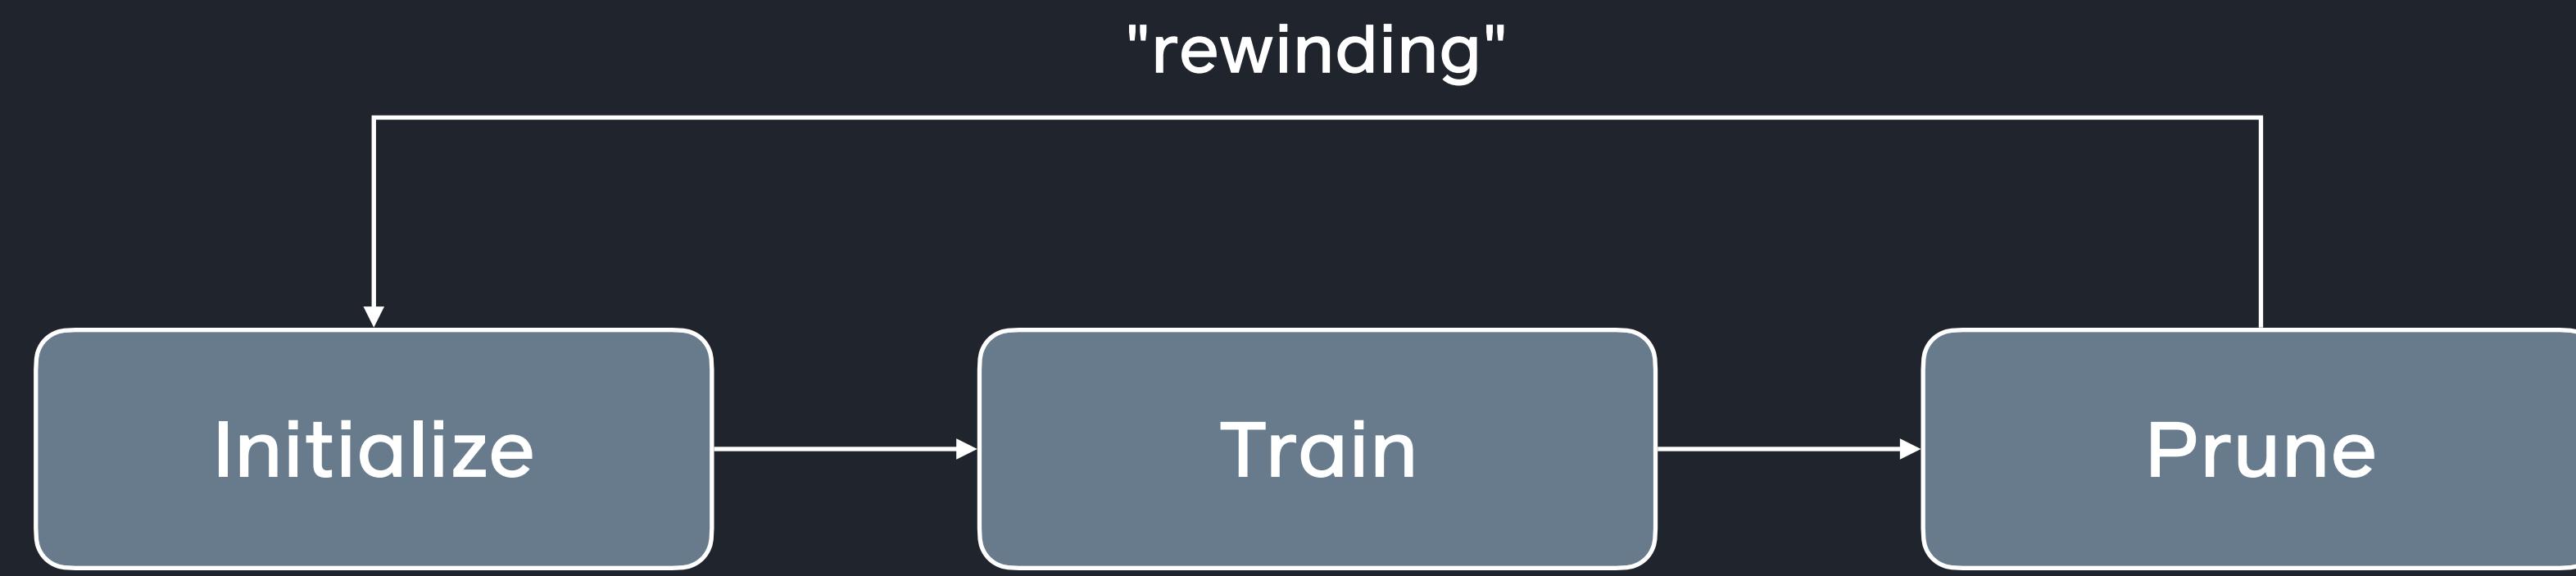
Section title



Section title









Section title

035bb6eaec977e3de60e0853d63cb011.bz2	4d2060ad068f9c8beb61c909ccedf1fb.bz2	9fd11c13ff649e45082663585e705ce9.bz2
035bb6eaec977e3de60e0853d63cb011.slim.json	4d2060ad068f9c8beb61c909ccedf1fb.slim.json	9fd11c13ff649e45082663585e705ce9.slim.json
049d06b3f12d4485d2c9555eb4f27b24.bz2	50394e6dc4e478c3a9db0ab660937987.bz2	b420996715b37da7749ce42badbada10.bz2
049d06b3f12d4485d2c9555eb4f27b24.slim.json	50394e6dc4e478c3a9db0ab660937987.slim.json	b420996715b37da7749ce42badbada10.slim.json
06cfe230e2b9565ab2d59e864eab7045.bz2	59f8a498b7ead01b98767ba912c8bb2f.bz2	ba803121bfff54079d3fadaa6ff434c53.bz2
06cfe230e2b9565ab2d59e864eab7045.slim.json	59f8a498b7ead01b98767ba912c8bb2f.slim.json	ba803121bfff54079d3fadaa6ff434c53.slim.json
08291cdd4c0f12de74166b02d935eabe.bz2	5a3771b0392256a2d5625d71ce663e2b.bz2	bc2887dc4df8d4baab6952ec8ef26323.bz2
08291cdd4c0f12de74166b02d935eabe.slim.json	5a3771b0392256a2d5625d71ce663e2b.slim.json	bc2887dc4df8d4baab6952ec8ef26323.slim.json
086ae9abf2282a3aa26e3f769f1ef509.bz2	5a3c5fb365bcf3935d7fd89e2145582.bz2	bc8a78325b9290b49671ece38ad2885.bz2
086ae9abf2282a3aa26e3f769f1ef509.slim.json	5a3c5fb365bcf3935d7fd89e2145582.slim.json	bc8a78325b9290b49671ece38ad2885.slim.json
0a74d268096a65c2a1779414aadba3df.bz2	61b999ebb35c25c34033a2987c92be76.bz2	c53d29419b6cfe7ff959b2b63ea782e1.bz2
0a74d268096a65c2a1779414aadba3df.slim.json	61b999ebb35c25c34033a2987c92be76.slim.json	c53d29419b6cfe7ff959b2b63ea782e1.slim.json
20d92fb2765e4edb8a8f75adbb58696.bz2	440a835c5a513e111f1029917bz2	80fd73ad23cf5790d5170b0087af8fd.bz2
20d92fb2765e4edb8a8f75adbb58696.slim	440a835c5a513e111f1029917bz2	80fd73ad23cf5790d5170b0087af8fd.slim.json
29ef5352b5d4b460bdaeb15c40f6c8c.bz2	653345f4e5cd28cac62f01c2c.slim.json	f414a5f1e8e11322c4643c8b9c4.bz2
29ef5352b5d4b460bdaeb15c40f6c8c.slim.json	653345f4e5cd28cac62f01c2c.slim.json	f414a5f1e8e11322c4643c8b9c4.slim.json
30195510da3f77e3a12efcf1d41b63a.bz2	69dc21cf4ea444271a511f1029917.bz2	cb7aaaca40af5a918fd9cae80110.bz2
30195510da3f77e3a12efcf1d41b63a.slim.json	69dc21cf4ea444271a511f1029917.slim.json	cb7aaaca40af5a918fd9cae80110.slim.json
32537b8ea57608d5bf18d82b5b62d078.bz2	6b109f29bc047f37b1417bed1291e72f.bz2	cb79aa59933f51df2f577b1bb7fb429.bz2
32537b8ea57608d5bf18d82b5b62d078.slim.json	6b109f29bc047f37b1417bed1291e72f.slim.json	cb79aa59933f51df2f577b1bb7fb429.slim.json
344ae86e68def0bfb43bd3cfb8fbec13.bz2	6b6aba4fececabab06e11ca61e.bz2	cde95b2e20d54cf0666228ac61985182.bz2
344ae86e68def0bfb43bd3cfb8fbec13.slim.json	6b6aba4fececabab06e11ca61e.slim.json	cde95b2e20d54cf0666228ac61985182.slim.json
373ebdfa64468e1b3a23293e4d04680f.bz2	76ff392ef7b2748ba3784472f4884bf0.bz2	d5e50946d8d6fb49339b1b3abeca403a.bz2
373ebdfa64468e1b3a23293e4d04680f.slim.json	76ff392ef7b2748ba3784472f4884bf0.slim.json	d5e50946d8d6fb49339b1b3abeca403a.slim.json
386885205bcba9494dbebf580afee8b2.bz2	77bbe568a85d535408c0a3c8d945ed5.bz2	dd4dc4556488cfaffa3e044d1300d1c5.bz2
386885205bcba9494dbebf580afee8b2.slim.json	77bbe568a85d535408c0a3c8d945ed5.slim.json	dd4dc4556488cfaffa3e044d1300d1c5.slim.json
38ef0199861b8a1a5d05556d7fed1e52.bz2	79b8a91c6ed9a9c8b5ea6dea047d906a.bz2	e090eac0b64924c34acd90adb7300d64.bz2
38ef0199861b8a1a5d05556d7fed1e52.slim.json	79b8a91c6ed9a9c8b5ea6dea047d906a.slim.json	e090eac0b64924c34acd90adb7300d64.slim.json
3b4e8afc175fd03729a830fd136c5684.bz2	7a2d512e2e28c2beafc75214bf81a2cb.bz2	e6ac41a441f1a1edaaeaf0f583e921d4.bz2
3b4e8afc175fd03729a830fd136c5684.slim.json	7a2d512e2e28c2beafc75214bf81a2cb.slim.json	e6ac41a441f1a1edaaeaf0f583e921d4.slim.json
3cb027ad63585f8eda23639226302f63.bz2	8110e7ed753ac476cb080ec43cdd8291.bz2	ec50fc8bd8523301dc7a1cc38e3133e1.bz2
3cb027ad63585f8eda23639226302f63.slim.json	8110e7ed753ac476cb080ec43cdd8291.slim.json	ec50fc8bd8523301dc7a1cc38e3133e1.slim.json
42fa6cdfb64714c43c3419381a8fb821.bz2		
42fa6cdfb64714c43c3419381a8fb821.slim.json		

Reproducible Experiment Orchestration

[facebookresearch/dagger](https://facebookresearch.github.io/dagger/)

dagger is a minimal framework for describing trees of network-mutating actions suited to the needs of researchers, allowing fast experimentation as well as maintenance of clear provenance in experiment evolution .

Goals:

- Allow researchers to abstract away *fundamental scientific contributions* from *experiment-tracking boilerplate code*
- Bookkeeping: track model state provenance

Concepts:

- **Experiment**: the graph
- **Experiment State**: a node
- **Recipe**: an edge

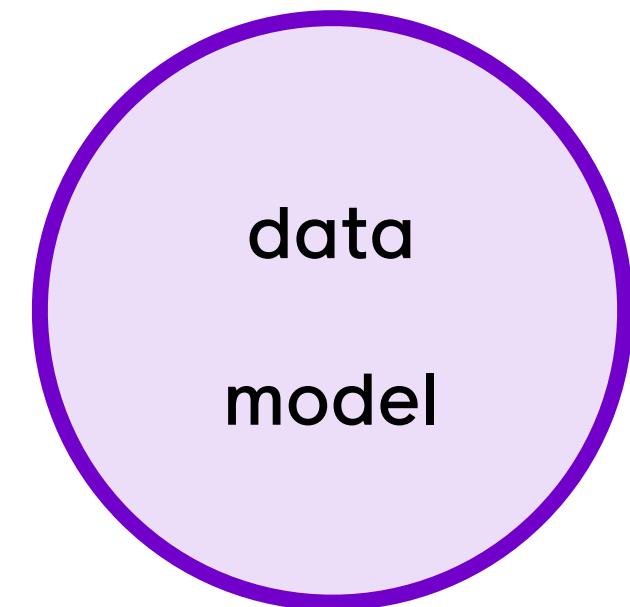
dagger is a minimal framework for describing trees of network-mutating actions suited to the needs of researchers, allowing fast experimentation as well as maintenance of clear provenance in experiment evolution .

Goals:

- Allow researchers to abstract away *fundamental scientific contributions* from *experiment-tracking boilerplate code*
- Bookkeeping: track model state provenance

Concepts:

- **Experiment**: the graph
- **Experiment State**: a node
- **Recipe**: an edge



Experiment State A

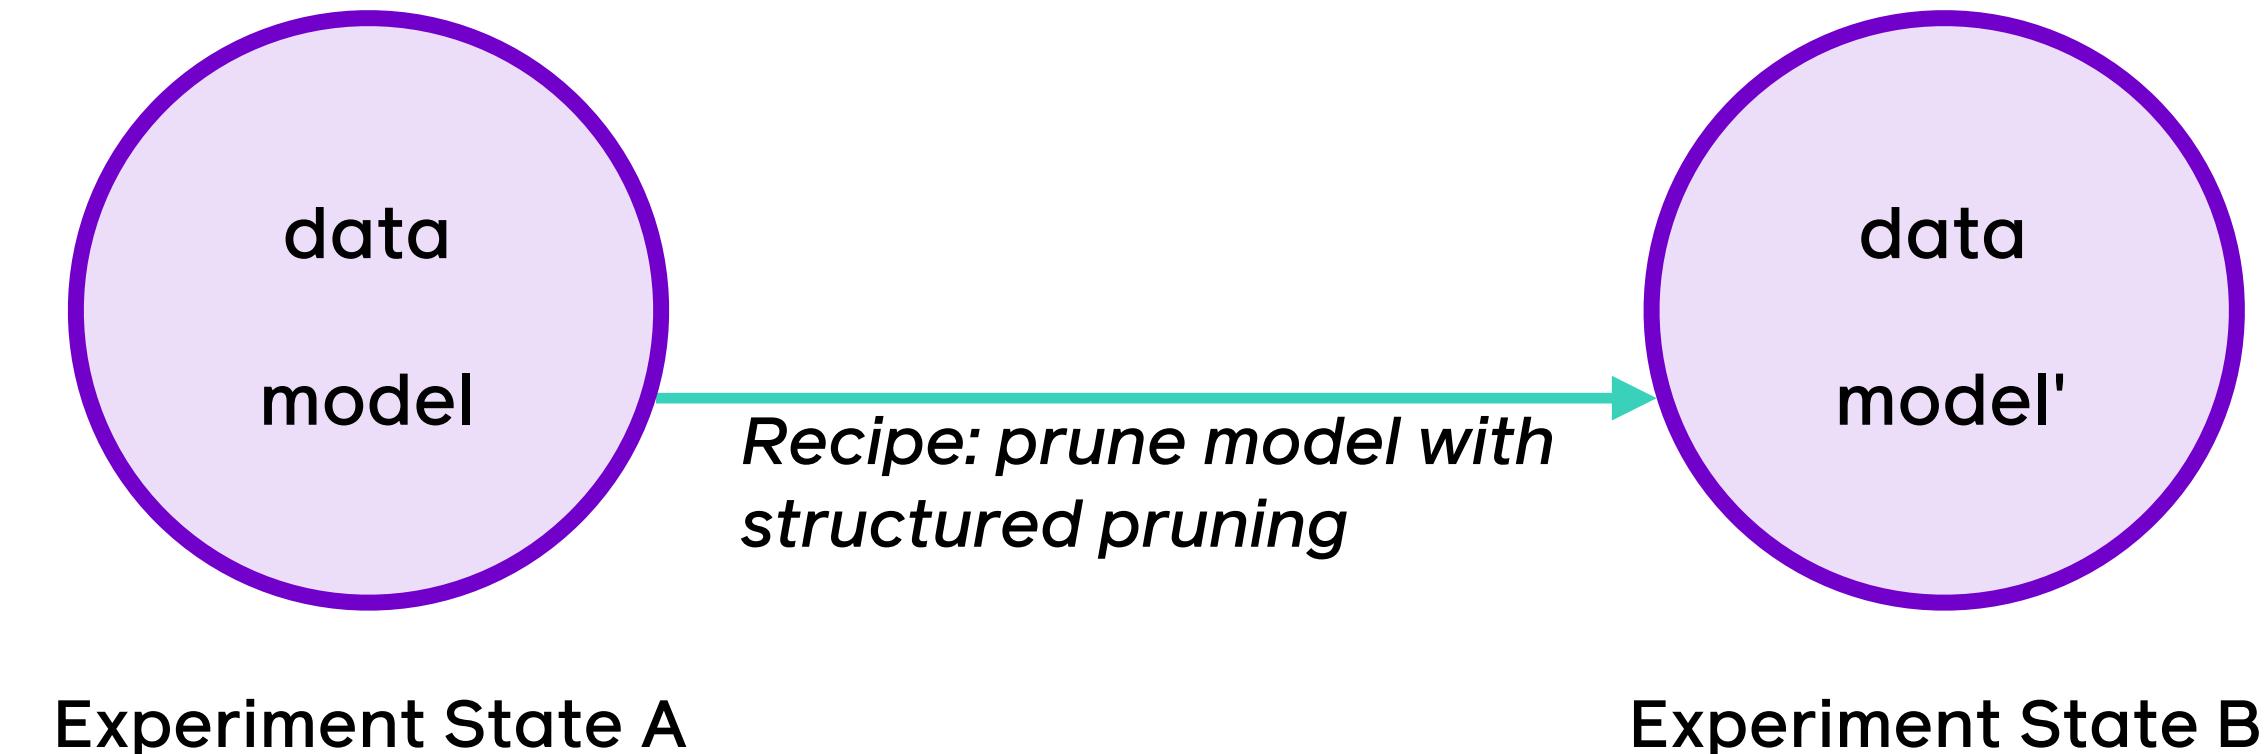
dagger is a minimal framework for describing trees of network-mutating actions suited to the needs of researchers, allowing fast experimentation as well as maintenance of clear provenance in experiment evolution .

Goals:

- Allow researchers to abstract away *fundamental scientific contributions* from *experiment-tracking boilerplate code*
- Bookkeeping: track model state provenance

Concepts:

- **Experiment**: the graph
- **Experiment State**: a node
- **Recipe**: an edge



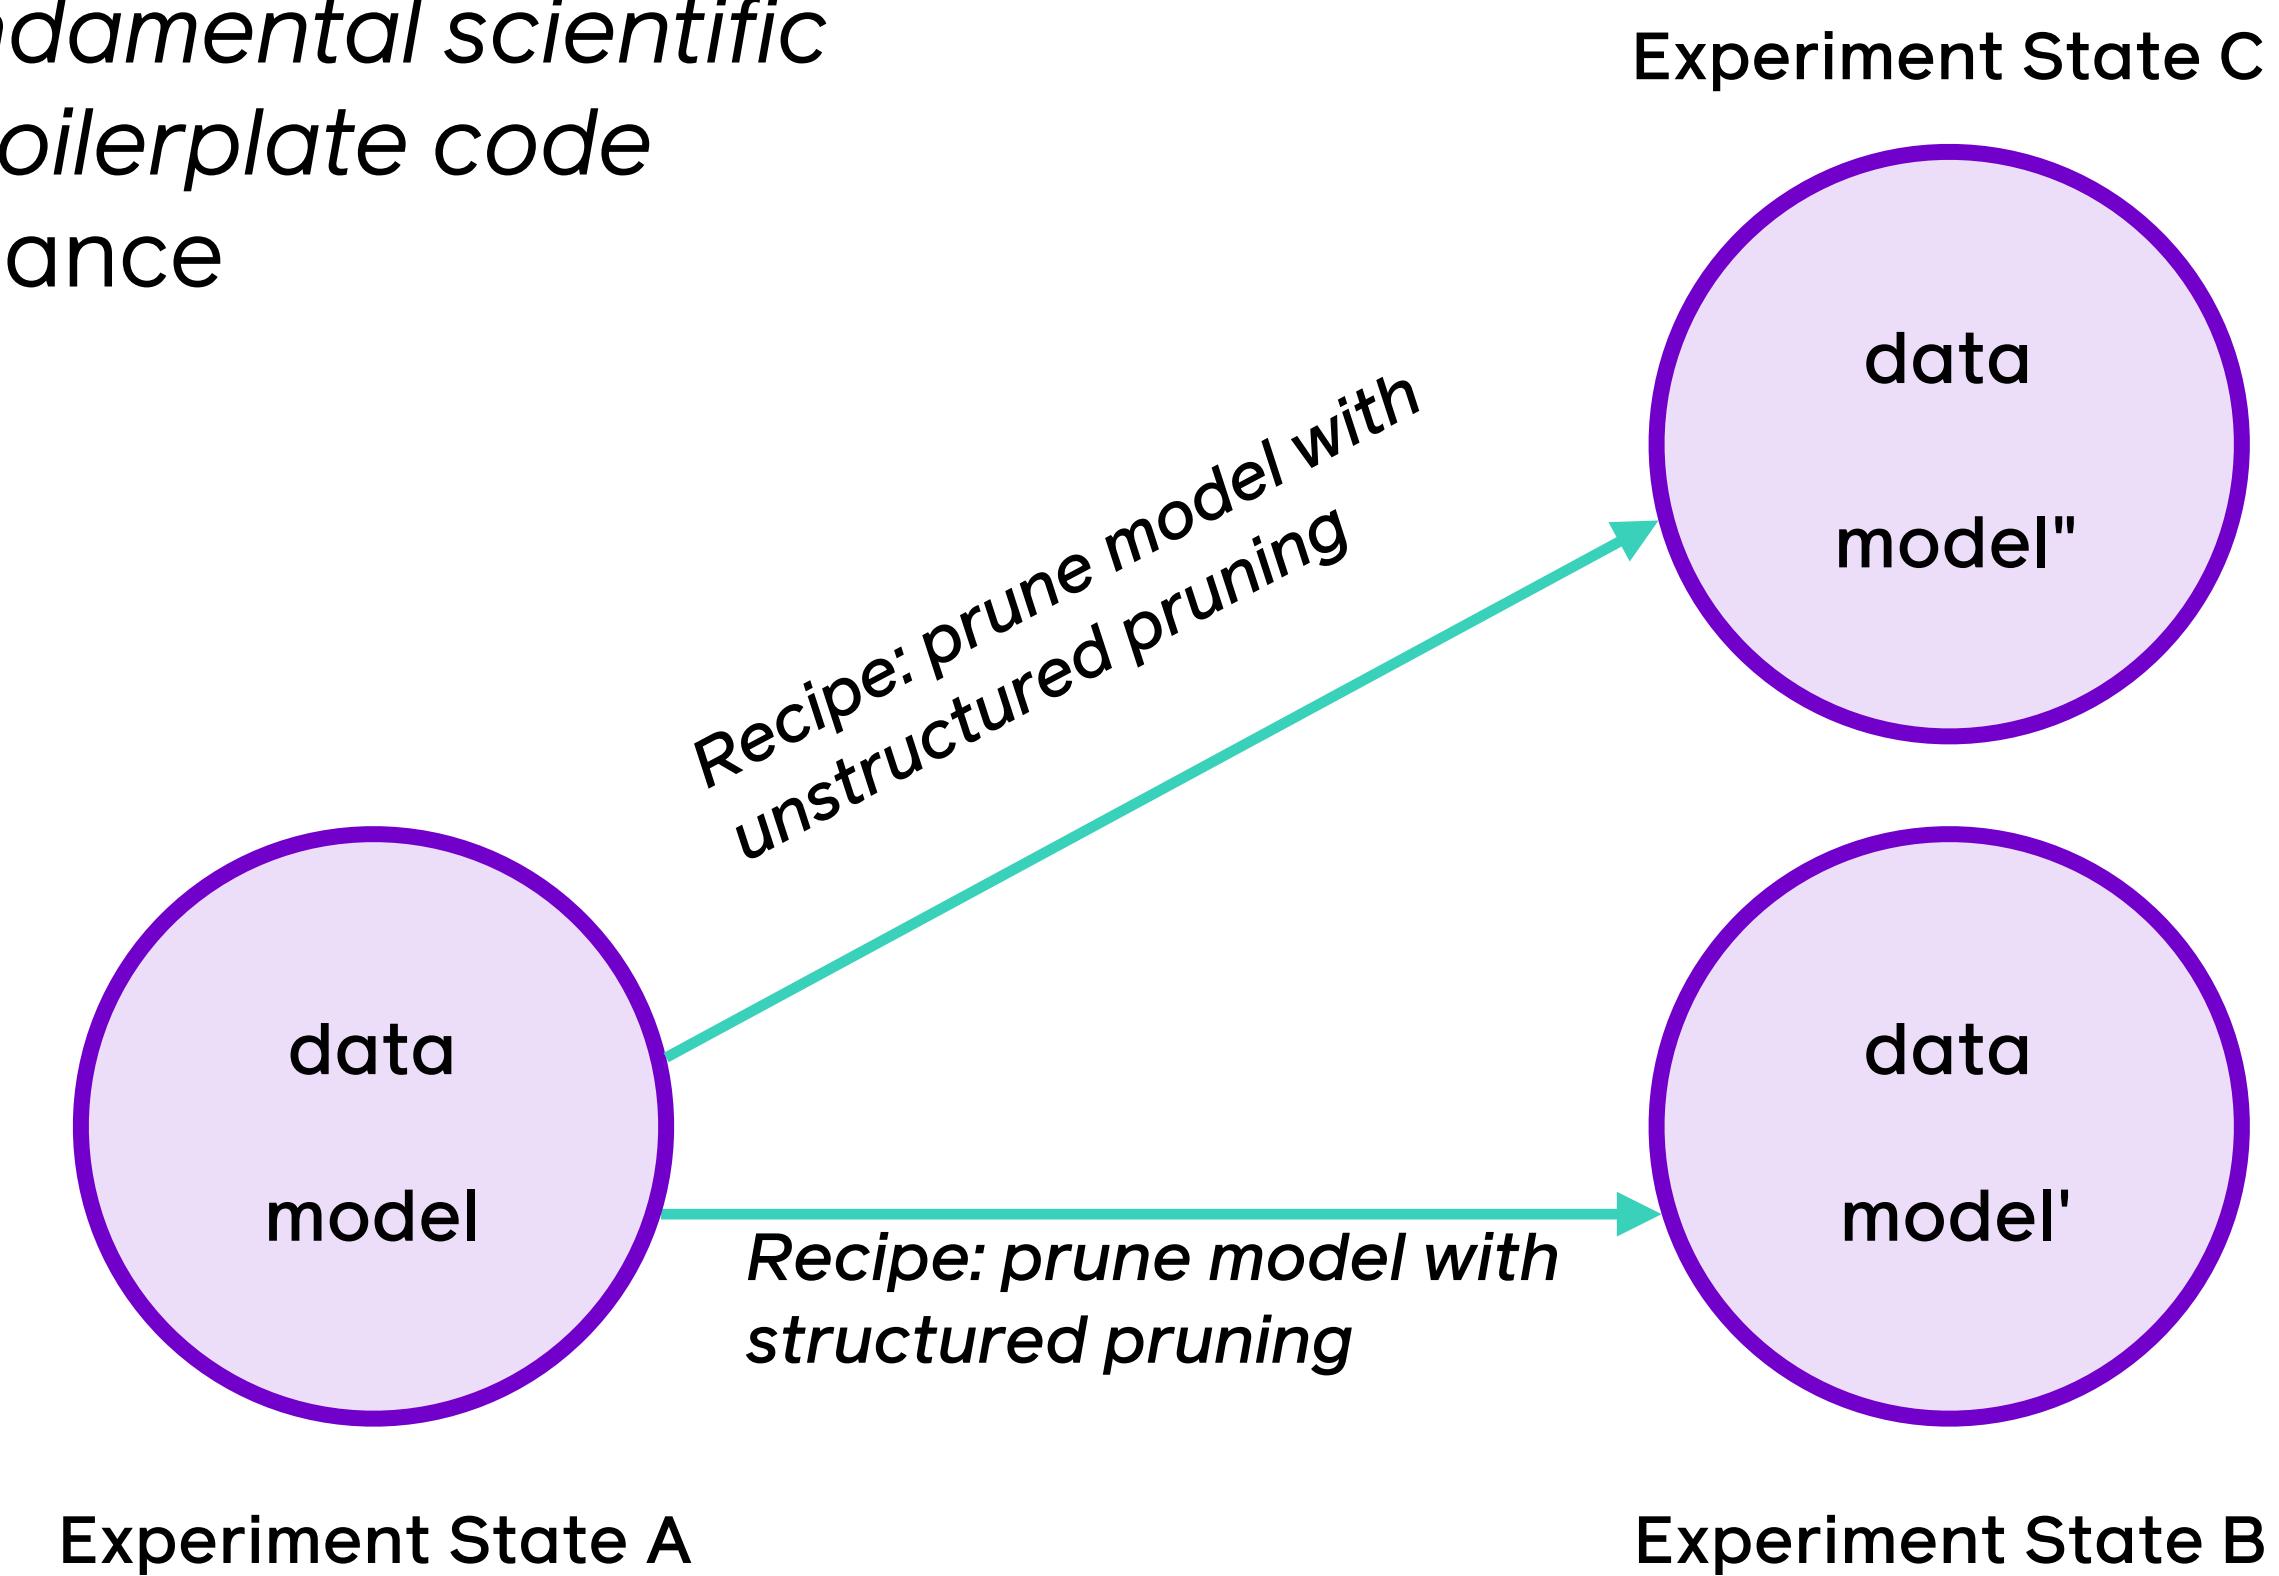
dagger is a minimal framework for describing trees of network-mutating actions suited to the needs of researchers, allowing fast experimentation as well as maintenance of clear provenance in experiment evolution .

Goals:

- Allow researchers to abstract away *fundamental scientific contributions* from *experiment-tracking boilerplate code*
- Bookkeeping: track model state provenance

Concepts:

- **Experiment**: the graph
- **Experiment State**: a node
- **Recipe**: an edge



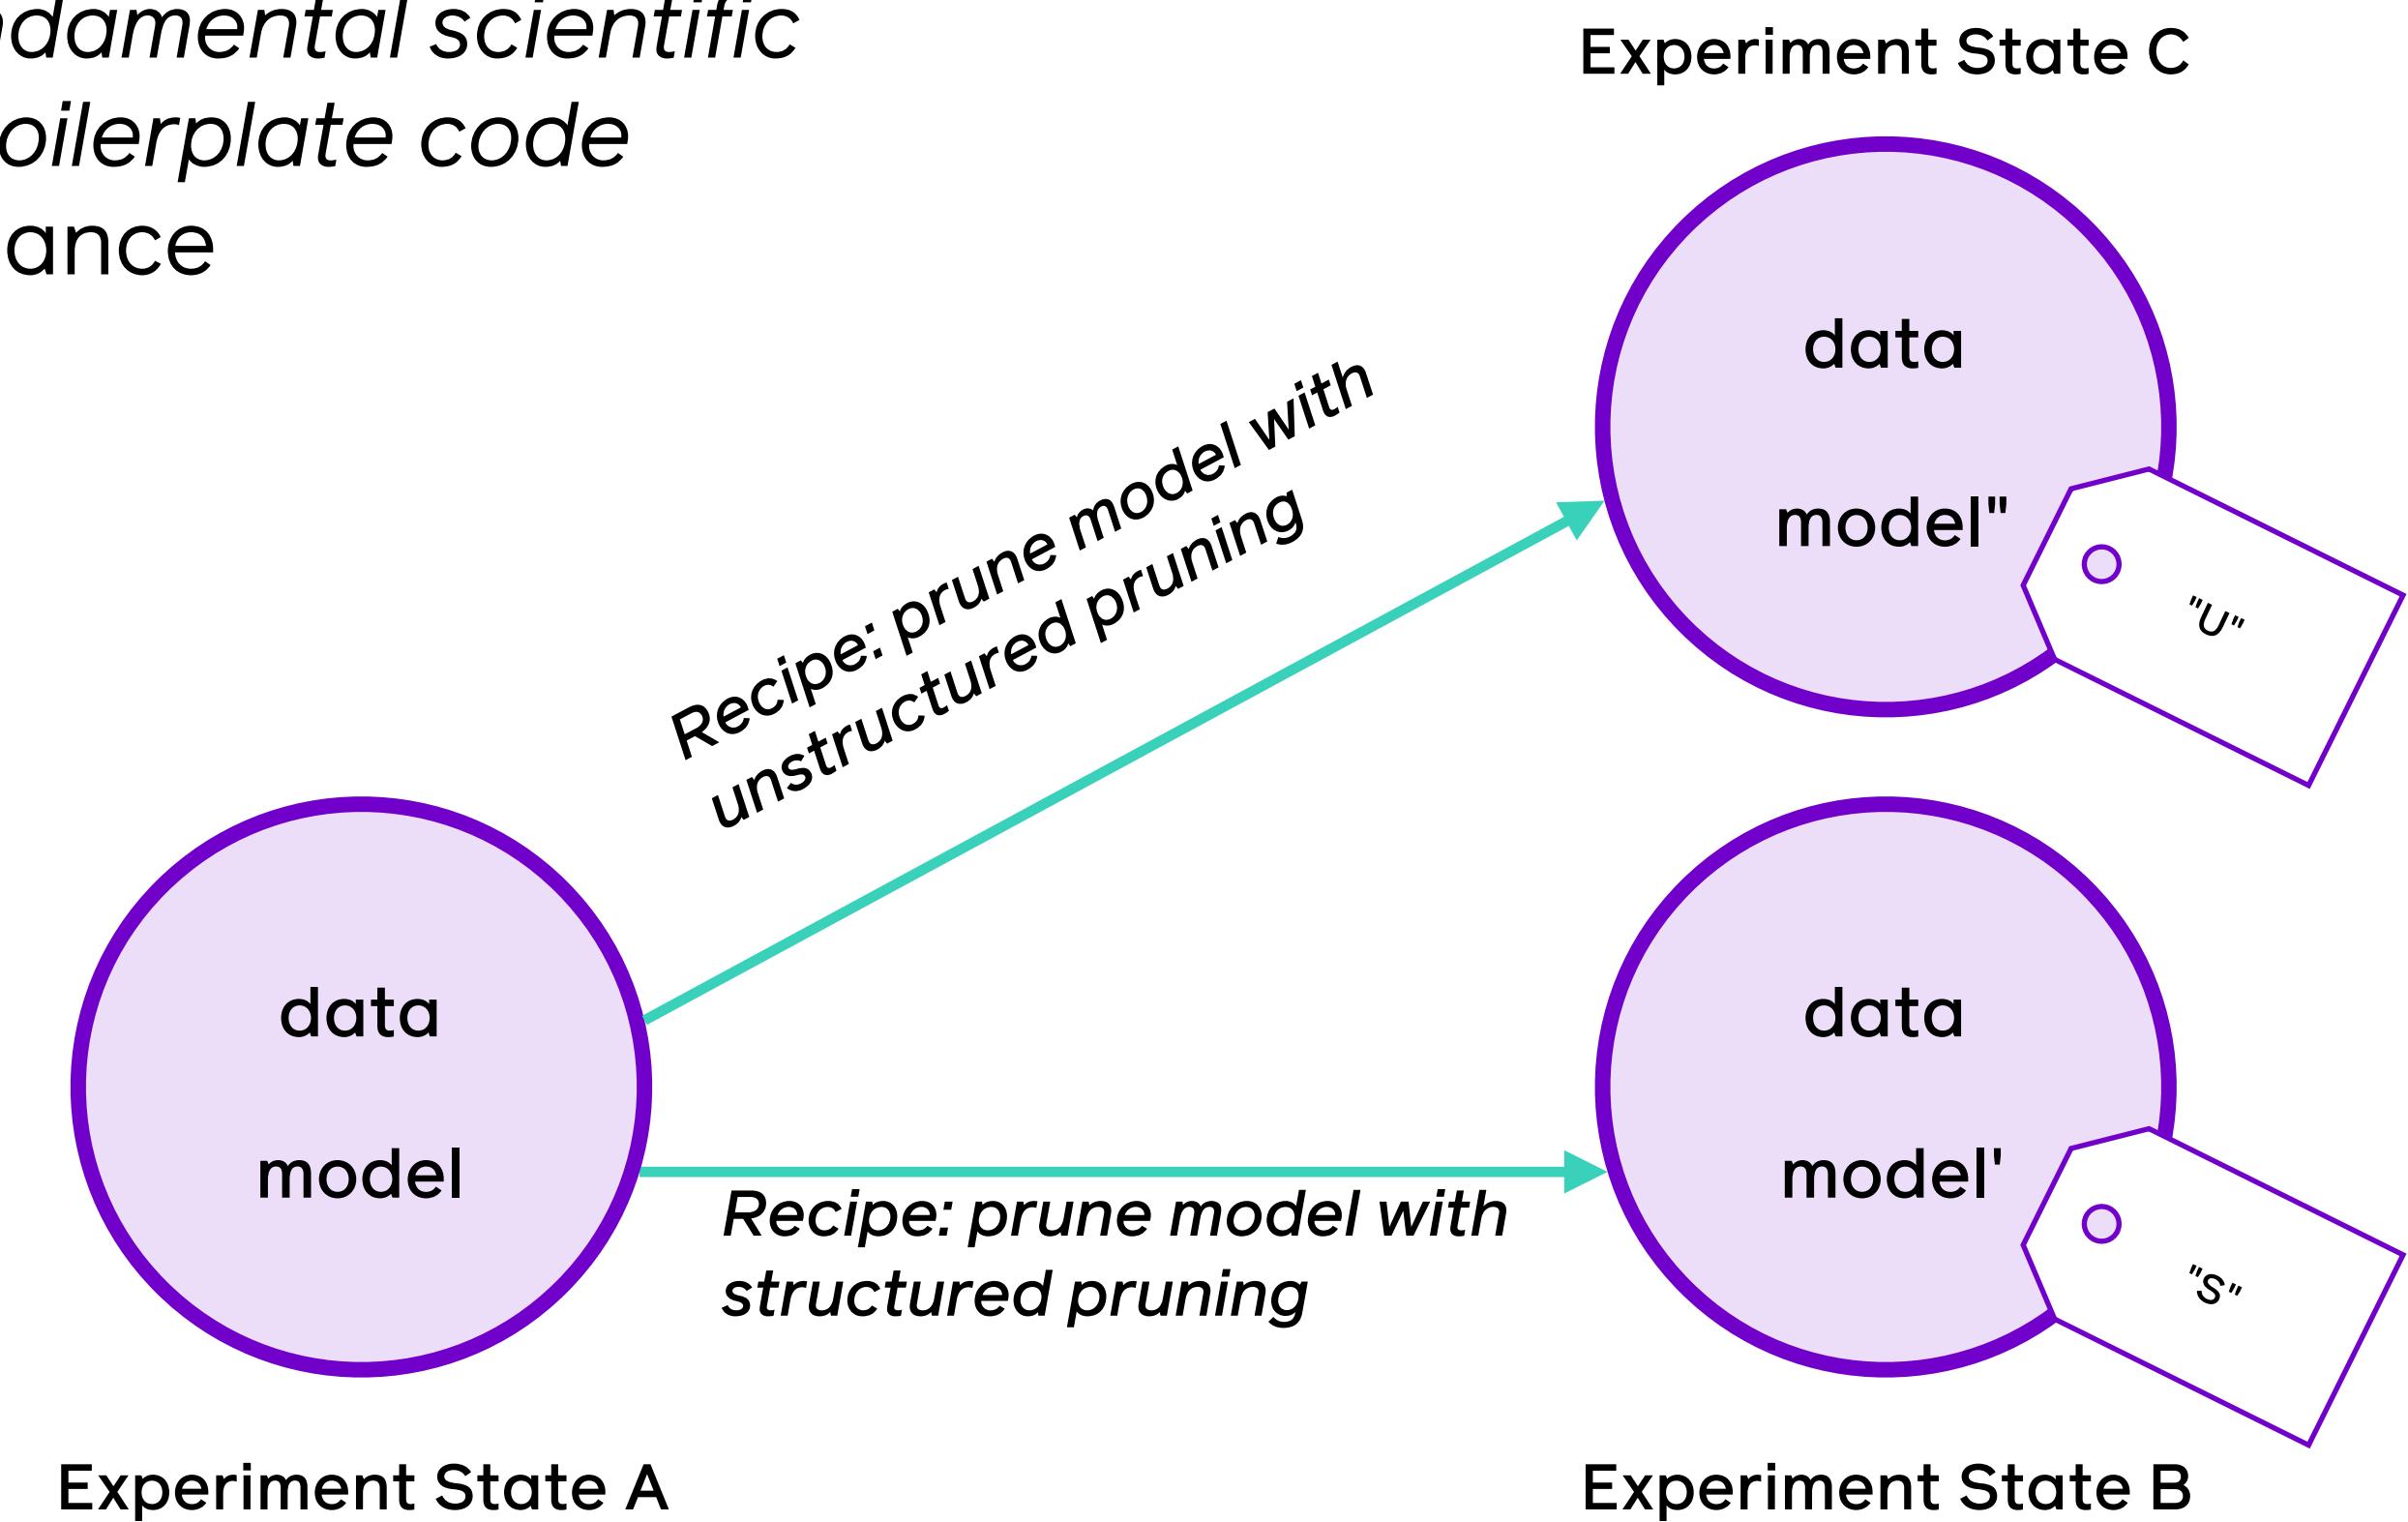
dagger is a minimal framework for describing trees of network-mutating actions suited to the needs of researchers, allowing fast experimentation as well as maintenance of clear provenance in experiment evolution .

Goals:

- Allow researchers to abstract away *fundamental scientific contributions* from *experiment-tracking boilerplate code*
- Bookkeeping: track model state provenance

Concepts:

- **Experiment**: the graph
- **Experiment State**: a node
- **Recipe**: an edge



Experiment Loop

```
1 exp = dg.Experiment("/path/to/experiment/folder", state_class=State)
2 root_state = exp.spawn_new_tree(dataset_name="cifar-10", model_name="vgg-11")
3
4 for lr in [0.01, 0.1]:
5     train = TrainRecipe(nb_epochs=100, lr=lr)
6     prune = PruneRecipe(pruning_technique="lowest_magnitude", pruning_fraction=0.2)
7     s = root_state
8     with exp.tag(f"lr:{lr}"):
9         s = train(s)
10        eval_fn(s)
11        with exp.tag("pruned"):
12            s = prune(s)
13 exp.run()
```

Experiment Analysis

```
1 >>> exp = Experiment.restore("/path/to/experiment/folder", slim=True)
2 >>> exp.graph.draw() # Draws the graph in Figure 1
3 >>> s = exp.graph.nodes.filter("pruned") & exp.graph.nodes.filter("lr:0.1")
4 >>> s[0].restore()
```

Custom Definitions

```
1 import dagger as dg
2 from yourlib import get_data, get_model, train_model, prune_model, eval_model
3
4 class State(dg.ExperimentState):
5
6     PROPERTIES = ["dataset_name", "model_name"]
7     NONHASHED_ATTRIBUTES = ["train_data", "eval_data", "model"]
8
9     def initialize_state(self, **kwargs):
10         self.train_data, self.eval_data = get_data(self.dataset_name)
11         self.model = get_model(self.model_name)
12
13 class TrainRecipe(dg.Recipe):
14
15     PROPERTIES = ["nb_epochs", "lr"]
16
17     def run(self, state):
18         train_model(state.model, state.train_data, self.nb_epochs, self.lr)
19         return state
20
21 class PruneRecipe(dg.Recipe):
22
23     PROPERTIES = ["pruning_technique", "pruning_fraction"]
24
25     def run(self, state):
26         prune_model(state.model, self.pruning_technique, self.pruning_fraction)
27         return state
28
29 @dg.function
30 def eval_fn(state):
31     eval_acc = eval_model(state.model, state.eval_data)
32     print(f"Experiment: {state.tags}, Accuracy: {eval_acc}")
```

Centralized Pruning in PyTorch

`torch.nn.utils.prune`

torch.nn.utils.prune

Different tensor pruning techniques enabled under a unified framework

BasePruningMethod

CLASS `torch.nn.utils.prune.BasePruningMethod`

[SOURCE]

Abstract base class for creation of new pruning techniques.

CLASSMETHOD `apply(module, name, *args, **kwargs)`

[SOURCE]

`apply_mask(module)`

[SOURCE]

ABSTRACT `compute_mask(t, default_mask)`

[SOURCE]

`prune(t, default_mask=None)`

[SOURCE]

`remove(module)`

[SOURCE]

New pruning technique?

Just subclass BasePruningMethod and implement `compute_mask`!

PruningContainer

CLASS `torch.nn.utils.prune.PruningContainer(*args)`

[SOURCE]

Container holding a sequence of pruning methods for iterative pruning. Keeps track of the order in which pruning methods are applied and handles combining successive pruning calls.

Identity

CLASS `torch.nn.utils.prune.Identity`

[SOURCE]

Utility pruning method that does not prune any units but generates the pruning parametrization with a mask of ones.

RandomUnstructured

CLASS `torch.nn.utils.prune.RandomUnstructured(amount)`

[SOURCE]

Prune (currently unpruned) units in a tensor at random.

L1Unstructured

CLASS `torch.nn.utils.prune.L1Unstructured(amount)`

[SOURCE]

Prune (currently unpruned) units in a tensor by zeroing out the ones with the lowest L1-norm.

RandomStructured

CLASS `torch.nn.utils.prune.RandomStructured(amount, dim=-1)`

[SOURCE]

Prune entire (currently unpruned) channels in a tensor at random.

LnStructured

CLASS `torch.nn.utils.prune.LnStructured(amount, n, dim=-1)`

[SOURCE]

Prune entire (currently unpruned) channels in a tensor based on their Ln-norm.

CustomFromMask

CLASS `torch.nn.utils.prune.CustomFromMask(mask)`

[SOURCE]

torch.nn.utils.prune

BasePruningMethod

CLASS `torch.nn.utils.prune.BasePruningMethod`

[SOURCE]

Abstract base class for creation of new pruning techniques.

CLASSMETHOD `apply(module, name, *args, **kwargs)`

[SOURCE]

`apply_mask(module)`

[SOURCE]

ABSTRACT `compute_mask(t, default_mask)`

[SOURCE]

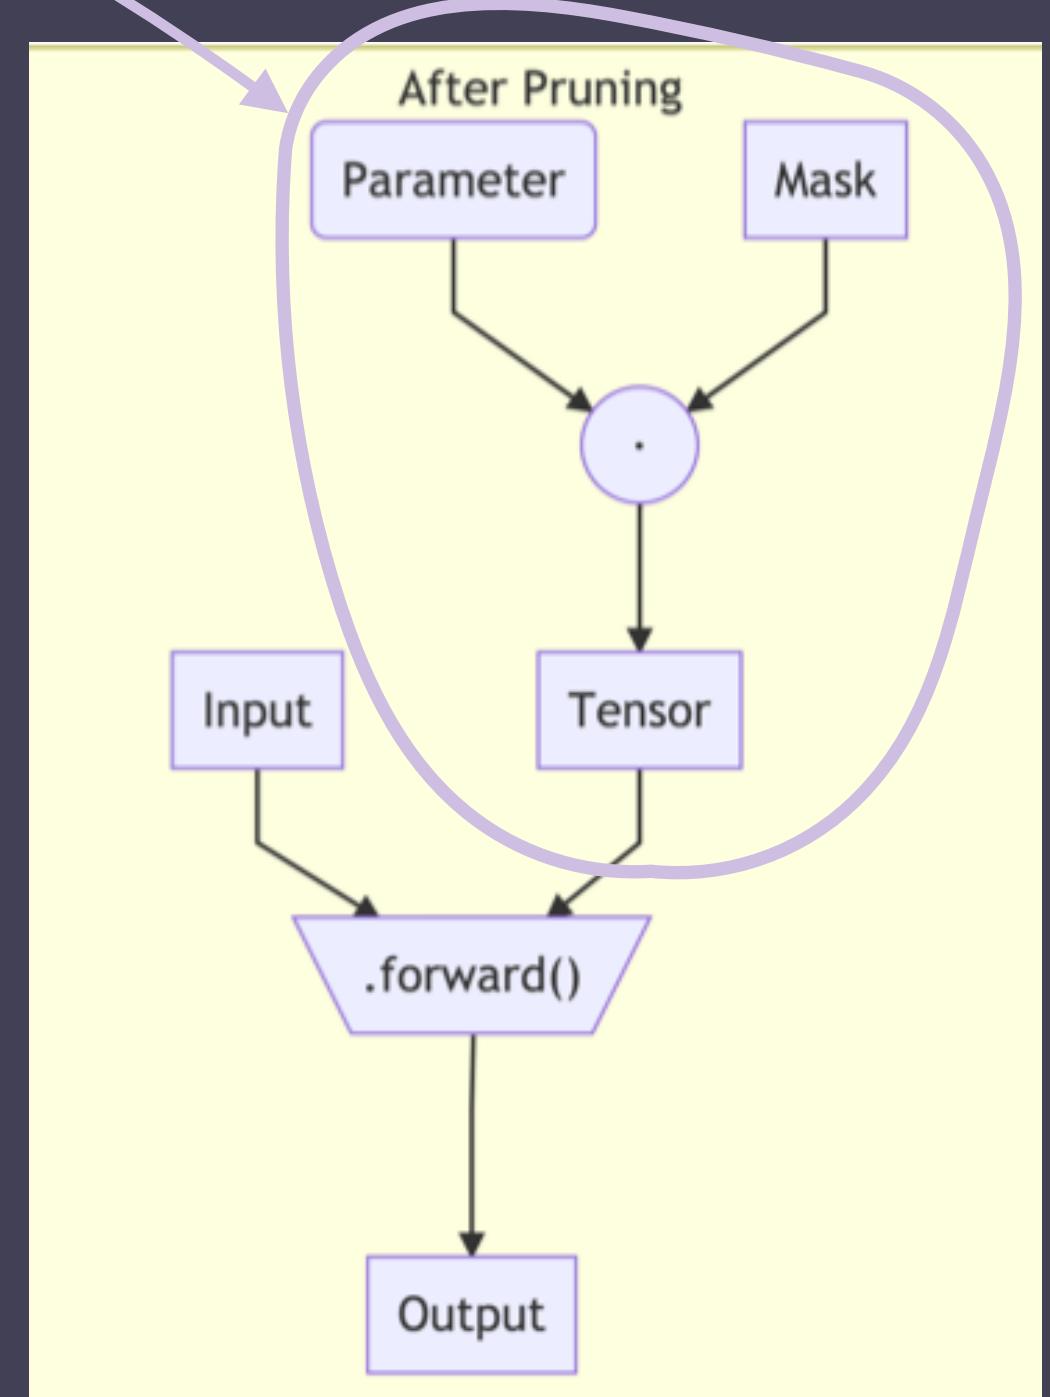
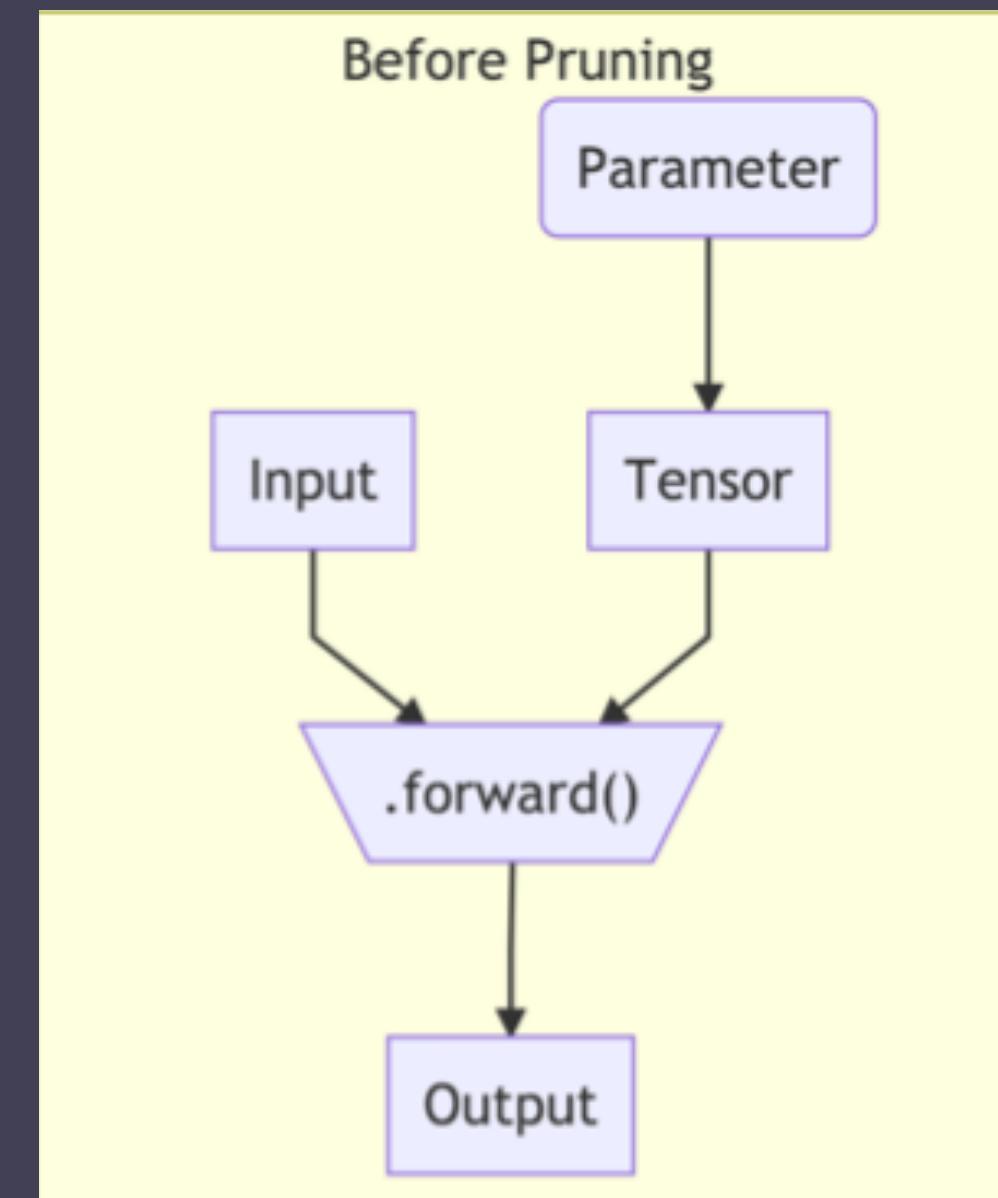
`prune(t, default_mask=None)`

[SOURCE]

`remove(module)`

[SOURCE]

Fetches the mask and the original, unpruned tensor to compute the pruned tensor during the forward pass → op is accounted for in the backward pass, too



torch.nn.utils.prune

BasePruningMethod

CLASS `torch.nn.utils.prune.BasePruningMethod`

[SOURCE]

Abstract base class for creation of new pruning techniques.

CLASSMETHOD `apply(module, name, *args, **kwargs)`

[SOURCE]

`apply_mask(module)`

[SOURCE]

ABSTRACT `compute_mask(t, default_mask)`

[SOURCE]

`prune(t, default_mask=None)`

[SOURCE]

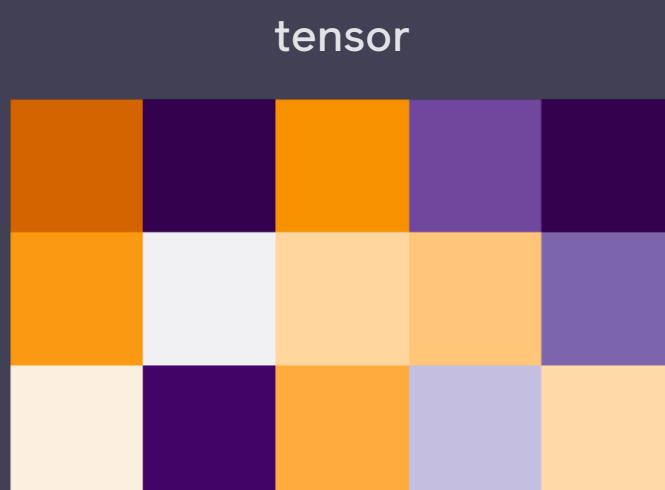
`remove(module)`

[SOURCE]

defines the interface → concrete subclasses must implement the logic

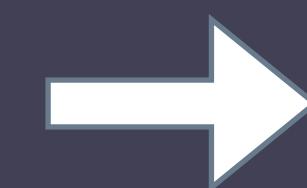
For example, in `prune.L1Unstructured`:

implements the logic that defines which portions of the tensors will be zeroed out while accounting for previously pruned entries



input

"remove lowest magnitude weights"



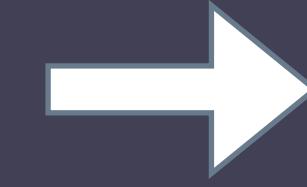
output

(through a `prune.PruningContainer`) it handles the case in which the tensor had previously been pruned by computing the valid entries in the tensor that can still be pruned and then applying the new pruning technique exclusively on those entries



previous mask

"remove lowest magnitude remaining weights"



input

output

Easy to use

```
model = LeNet() # unpruned model

# L_2 structured pruning will remove 50% of channels across axis 0
prune.ln_structured(
    module=model.conv1,
    name="weight",
    amount=0.5,
    n=2,
    dim=0
)
```

Iterative pruning made easy

prune.PruningContainer handles the combination of successive masks for you

```
for _ in range(10):
    # Remove 2 connections per iteration
    prune.l1_unstructured(module=model.fc1, name="bias", amount=2)
```

Global pruning made easy

```
parameters_to_prune = (
    (model.conv1, "weight"),
    (model.conv2, "weight"),
    (model.fc1, "weight"),
)

prune.global_unstructured(
    parameters_to_prune,
    pruning_method=prune.L1Unstructured,
    amount=0.2,
)
```

Easy to extend

```
class FooBarPruningMethod(prune.BasePruningMethod):
    """Prune every other entry in a tensor
    """

    PRUNING_TYPE = 'unstructured'

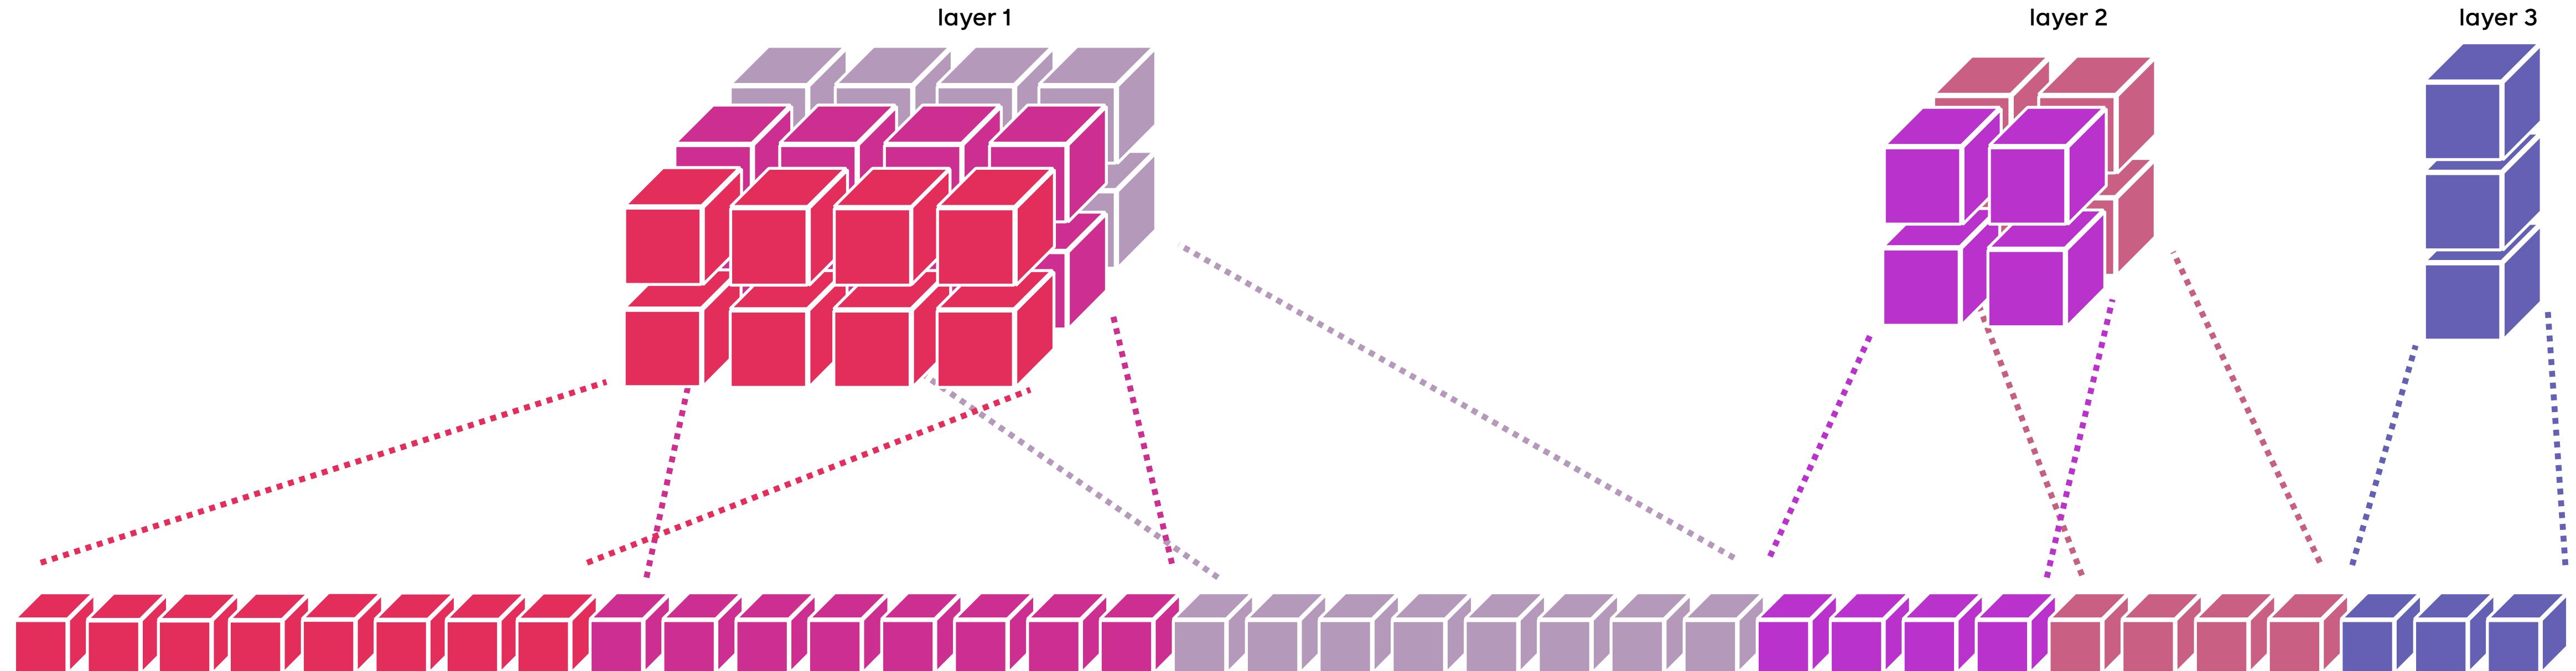
    def compute_mask(self, t, default_mask):
        mask = default_mask.clone()
        mask.view(-1)[::2] = 0
        return mask

    def foobar_unstructured(module, name):
        FooBarPruningMethod.apply(module, name)
        return module
```

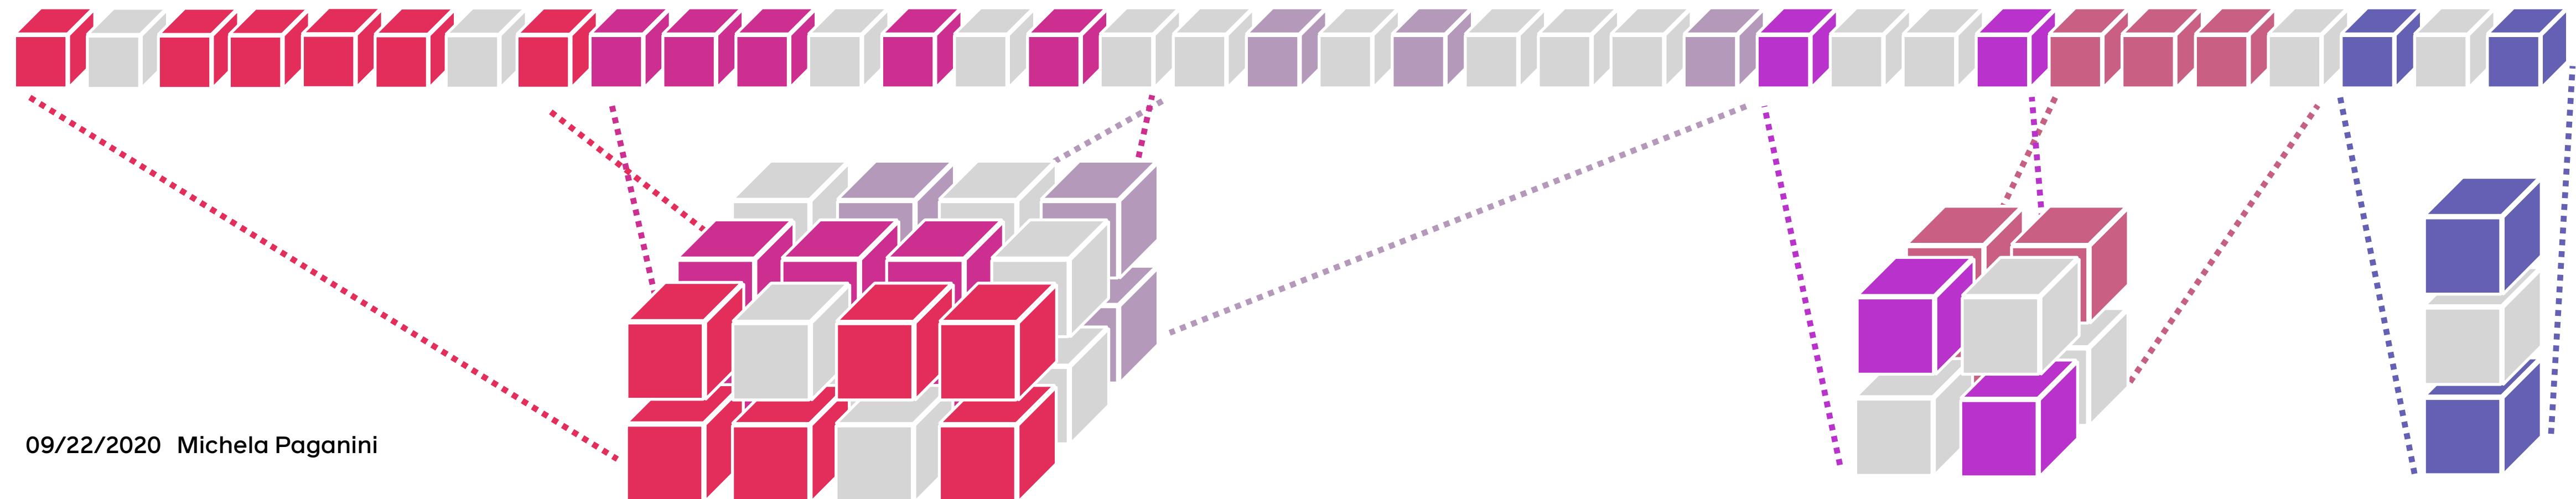
supports 3 PRUNING_TYPES:
'global', 'structured',
and 'unstructured' (to
determine how to combine
masks if pruning is applied
iteratively)

instructions on how to
compute the mask for the
given tensor according to
the logic of your pruning
technique

GlobalPruning



```
torch.nn.utils.prune.global_unstructured(...)
```



torch.nn.utils.prune

BasePruningMethod

CLASS `torch.nn.utils.prune.BasePruningMethod`

[SOURCE]

Abstract base class for creation of new pruning techniques.

CLASSEMTHOD `apply(module, name, *args, **kwargs)`

[SOURCE]

`apply_mask(module)`

[SOURCE]

ABSTRACT `compute_mask(t, default_mask)`

[SOURCE]

`prune(t, default_mask=None)`

[SOURCE]

`remove(module)`

[SOURCE]

`torch.nn.utils.prune` is designed to act on a `torch.nn.Module`

provides an interface for acting directly on a tensor

```
tensor = torch.randn([3, 5])
p = torch.nn.utils.prune.LnStructured(amount=1, dim=1, n=2)
masked_tensor = p.prune(tensor)
```

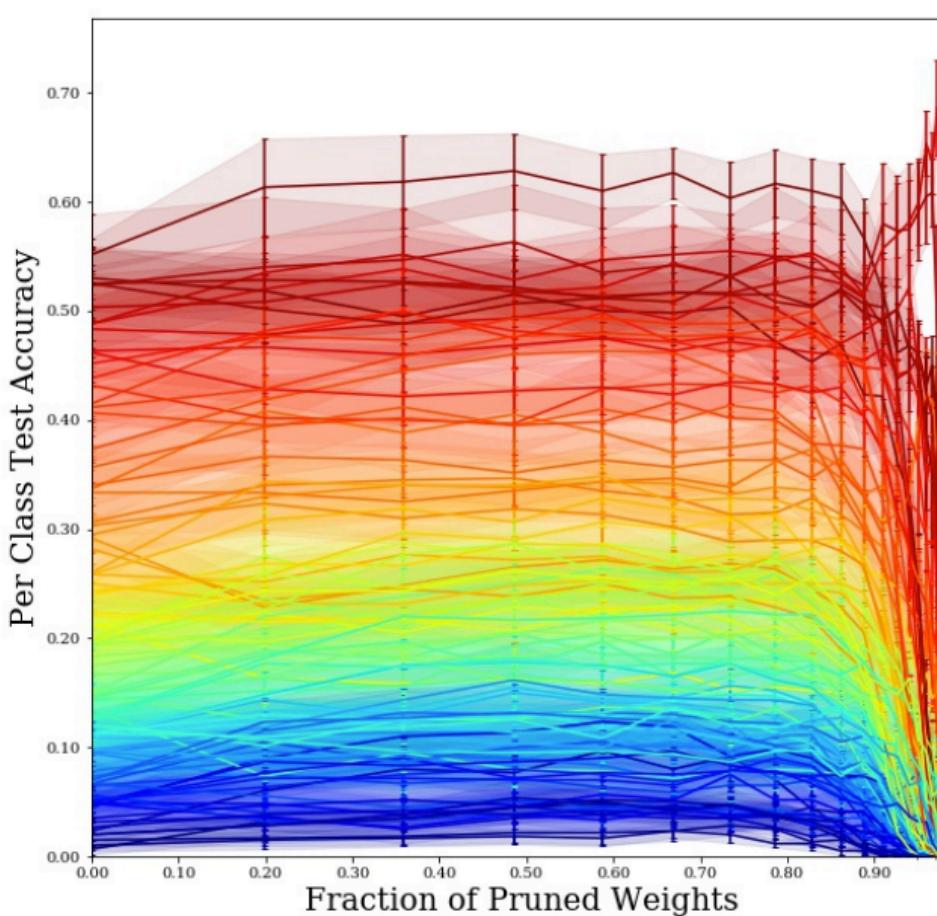
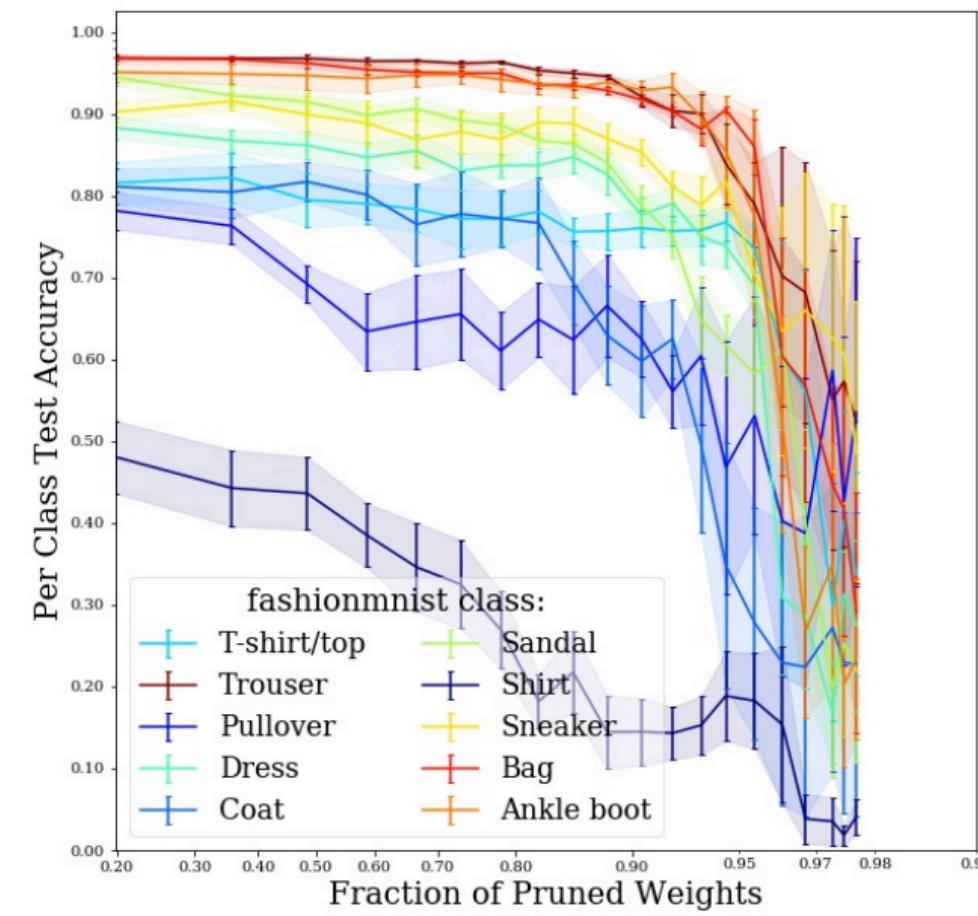
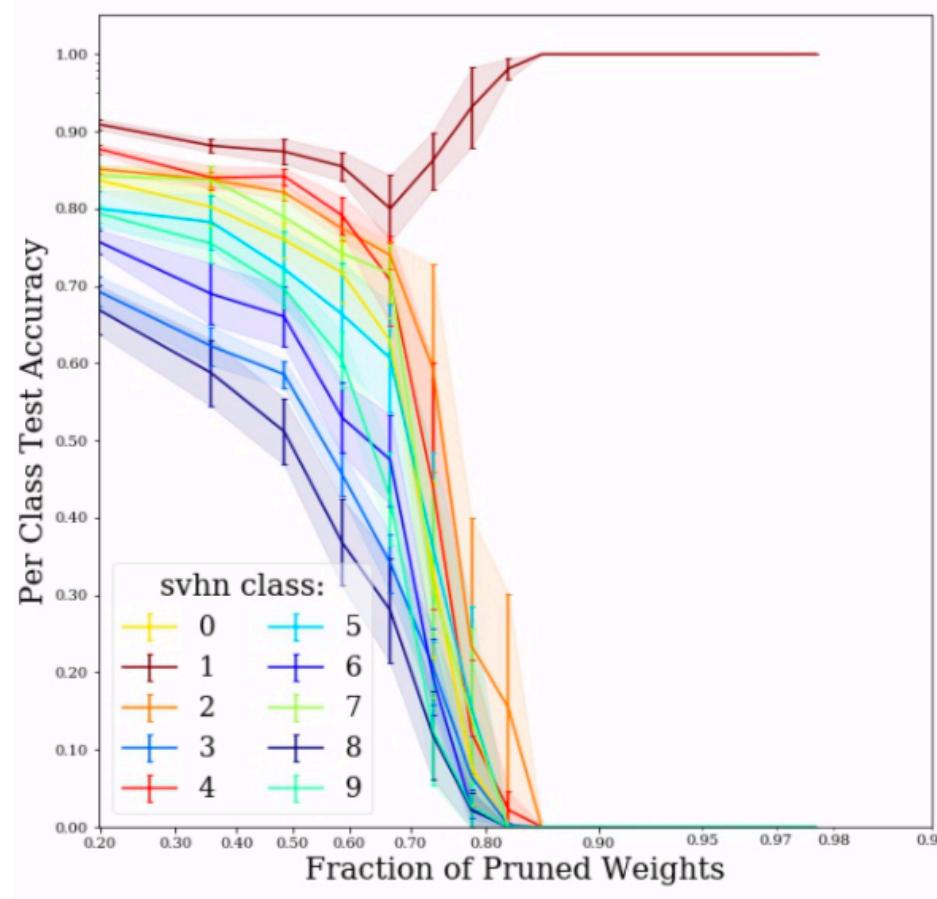
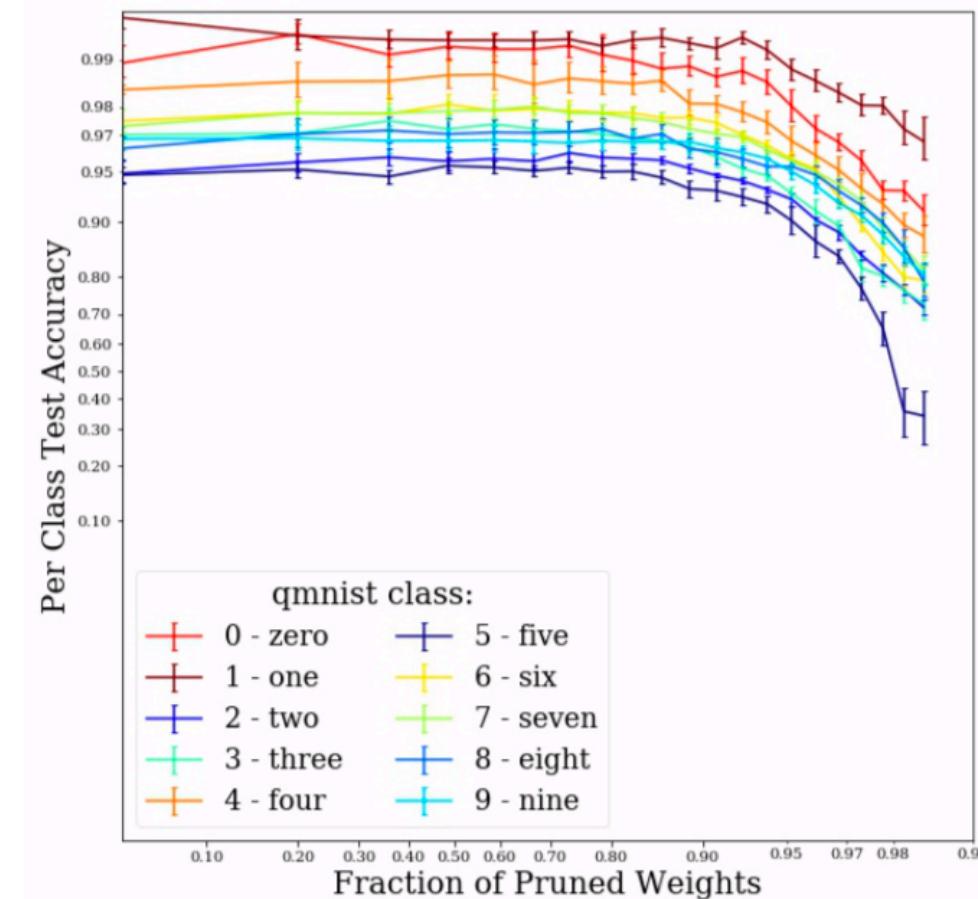
`torch.nn.utils.prune`



Prune Responsibly

arXiv:2009.09936

Test hypotheses that class complexity, difficulty, and representation matter in determining the accuracy after pruning



Prune and measure class accuracy for over 1M classes across over 100k models

- Fit a linear model for class accuracy as a function of:
- unpruned model class accuracy
 - class entropy
 - class representation
 - sparsity
 - dataset
 - model
 - pruning technique
 - weight treatment after pruning

Reject hypothesis that coefficients = 0

Closing Remarks

Papers with Code @paperswithcode · Sep 4

🔥 Introducing... the ML Reproducibility Challenge 2020! The 4th annual edition now expands to cover papers from 7 major ML conferences: NeurIPS, EMNLP, ACL, ICML, ICLR, CVPR and ECCV. Find more here at Papers With Code:

ML Reproducibility Challenge 2020

for papers published in:

Papers with Code - ML Reproducibility Challenge 2020
The ML Reproducibility Challenge 2020 covering paper published in seven major ML conferences: NeurIPS, ACL, EMNLP, ICLR, ICML, CV...
paperswithcode.com

Michela Paganini 🤝👩‍💻 ❤️ @WonderMicky · Sep 2

Upset at the latest [#NeurIPS2020](#) reviews? Interested in an alternative publication model for machine learning research? Looking for hypothesis-based science in AI? 🧠💡

Check out and submit to the [preregister.science](#) workshop at NeurIPS. **#SCIENCE!**
Deadline: Oct 7. More info ↗



NeurIPS2020 pre-registration workshop
Testing whether pre-registration can help fix our peer review system.
preregister.science

Thanks!

Questions? Contact me: michela@fb.com

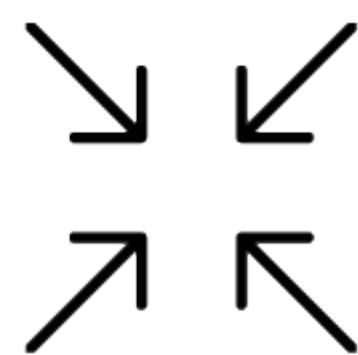
 WonderMicky

Learn from other Sciences.

Theoretical Science



Experimental Science



Engineering

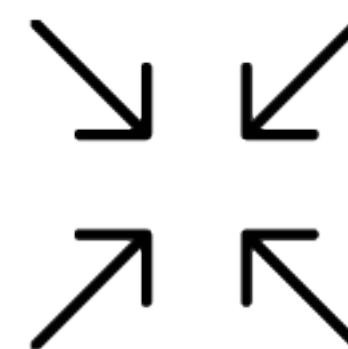
0101
1001
0110

Learn from other Sciences.

Theoretical Science



Experimental Science



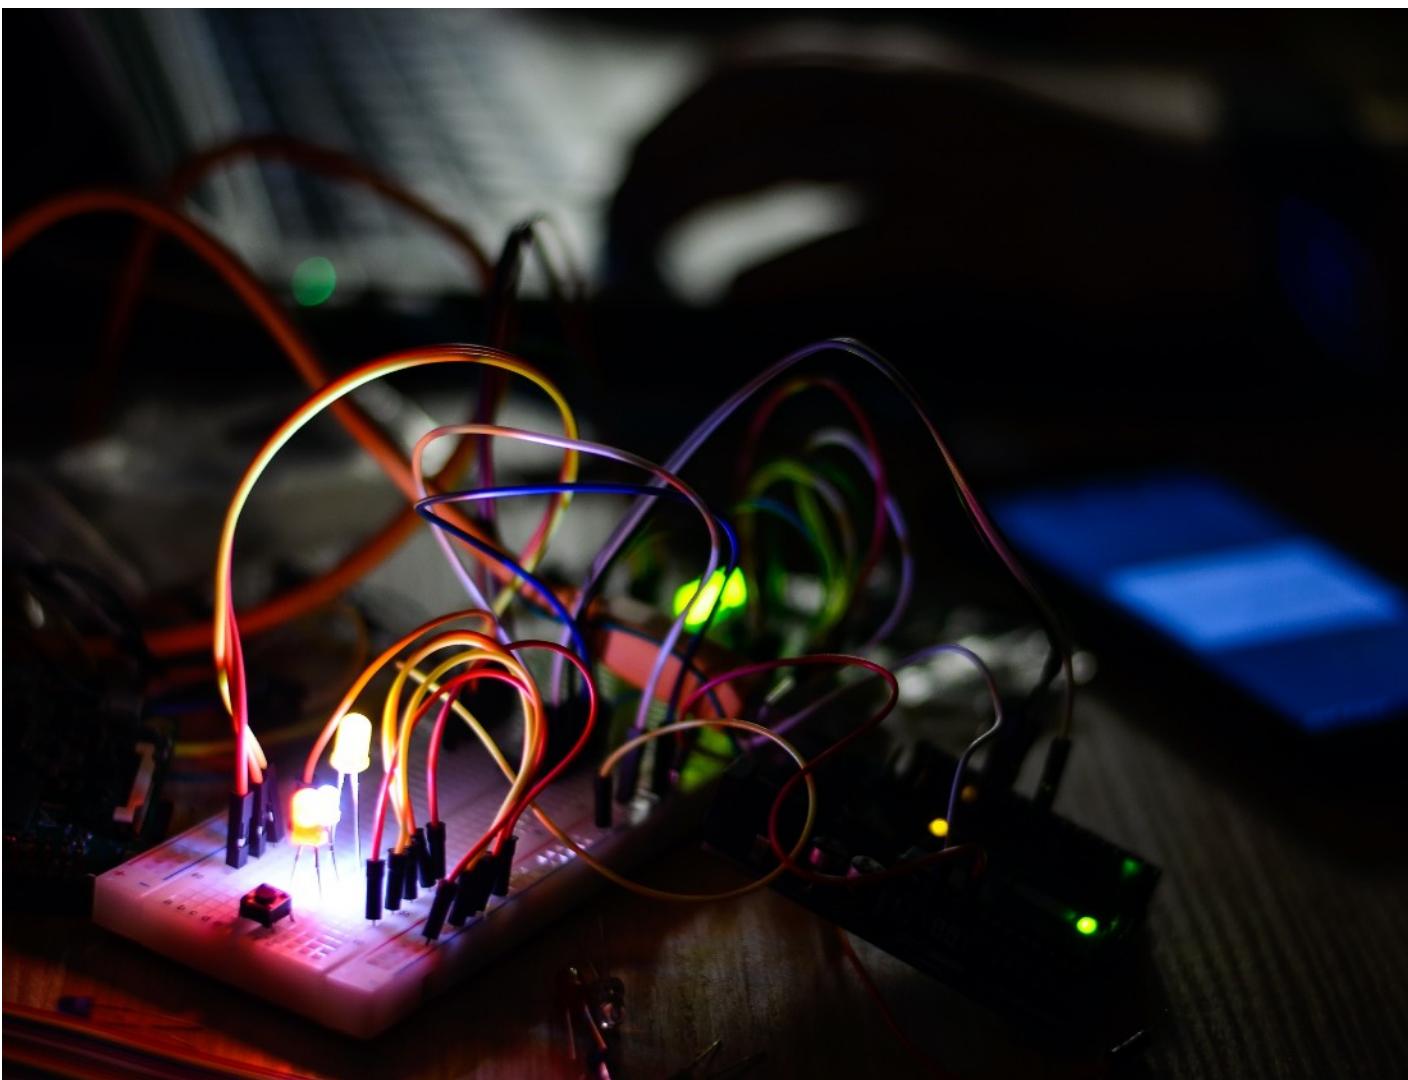
Engineering

0101
1001
0110

2. The scientific method in the science of ML

Neural Networks can be thought of as physical objects obeying laws of dynamics.

CAN STUDY THE INTERACTIONS OF THEIR FUNDAMENTAL COMPONENTS USING EXPERIMENTAL PROCEDURES.



$$\begin{cases} \frac{dS}{dt} = \frac{q_{act}}{q_{act}} - \beta(N-N_0)(1-\varepsilon S)S + \frac{N_e}{T_n} - \frac{N}{T_p} \\ \frac{dP}{dt} = T_b q_{po}(N-N_0)(1-\varepsilon S)S + \frac{\beta N}{T_n} - \frac{S}{T_p} \end{cases} \quad \left\{ \begin{array}{l} N=1 \\ P_f=(n) \end{array} \right.$$
$$S = \frac{T_p x_0}{T_p + T_n(1-\varepsilon)}$$
$$P_f = \frac{1}{1-\varepsilon}$$

"Grounding ML research in statistically sound hypothesis testing with careful control of nuisance parameters may encourage the publication of advances that stand the test of time."

Code Submission Policies

- ICML 2019 and NeurIPS 2019 rolled out explicit code-submission policies
- Many concerns regarding Dataset confidentiality, Proprietary software, Computation infrastructure, Replication of mistakes...
- NeurIPS 2019/2020 code submission policy leaves significant time and flexibility - “*expects code only for accepted papers, and only by the camera-ready deadline*”

Percentage of papers with code

